

29.03.2019

Statistical Methods in AI (CSE/ECE 471)

Lecture-20: ML for Sequential Data – Recurrent Neural Networks

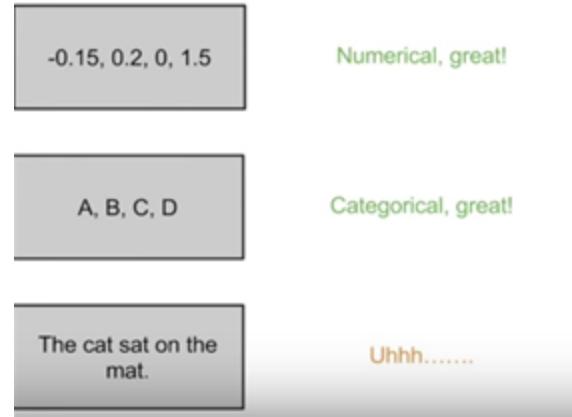
↳

Ravi Kiran

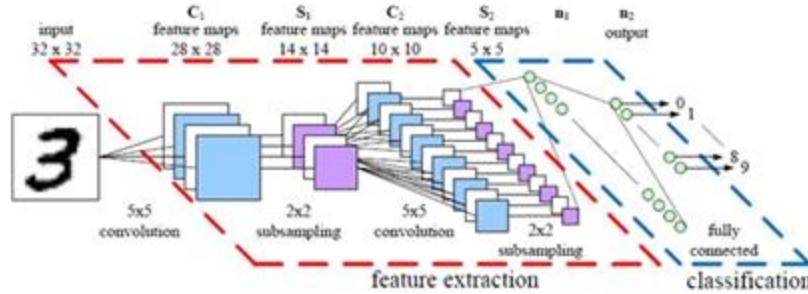
Center for Visual Information Technology (CVIT), IIIT Hyderabad



Data/Output Paradigms

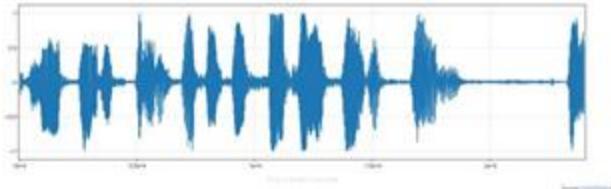


CNN Revisited

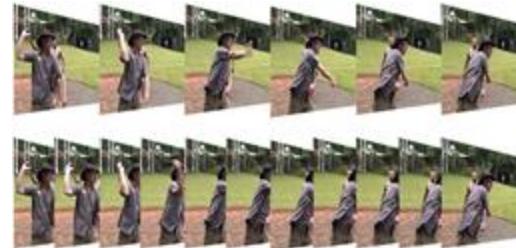


- Data is fed “in one shot”, not in parts
 - E.g. Image of “3” above is not processed row-by-row
- Captures “spatial context”

Sequential data



Audio

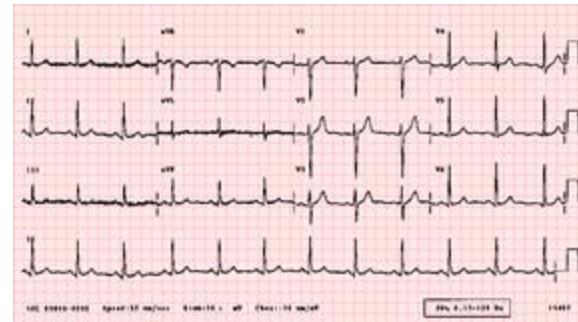


Video

O Nachiketa, after pondering well the pleasures that are or seem to be delightful, you have renounced them all. You have not taken the road abounding in wealth, where many men sink. (*Kathopanishad*, II:3)

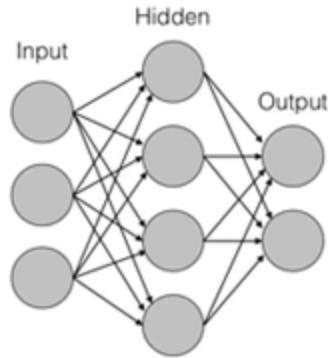
Text

many more

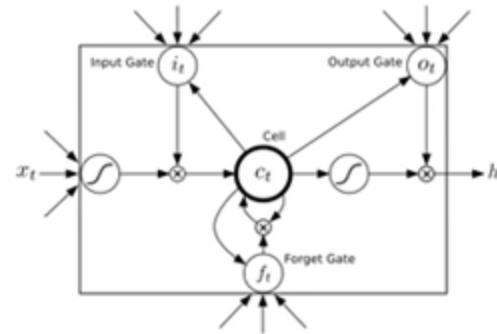


Biological

What are RNNs for ?



- Independence
- Fixed Length



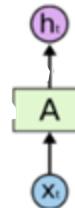
- Temporal dependencies
- Variable sequence length

Modelling a RNN

$$f(n) = f(n-1) + f(n-2)$$

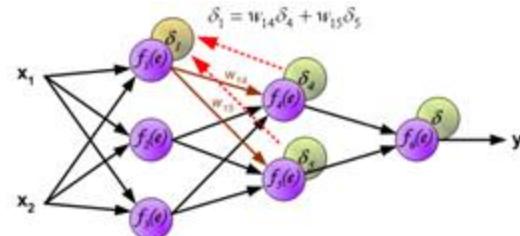
$$f(n-1) = f(n-2) + f(n-3)$$

:

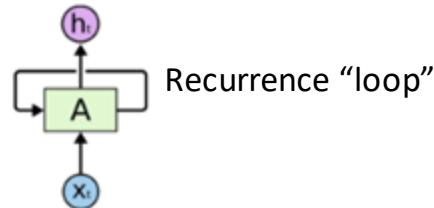


Error derivatives w.r.t weights in kth layer = f (Error derivatives from (k+1)th layer)

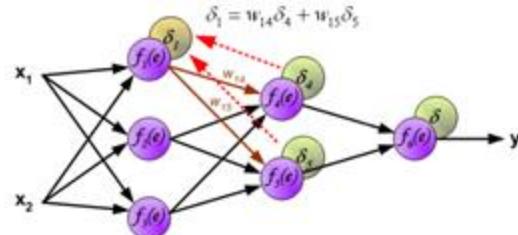
$$f(n) = g(f(n-k))$$



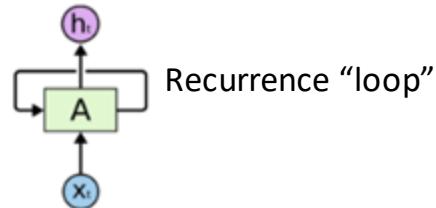
Modelling a RNN



Error derivatives w.r.t weights in kth layer = $f(\text{Error derivatives from (k+1)th layer})$

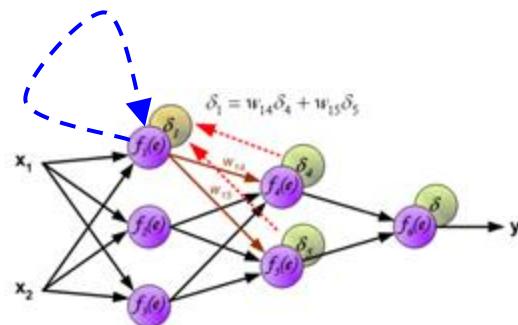


Modelling a RNN

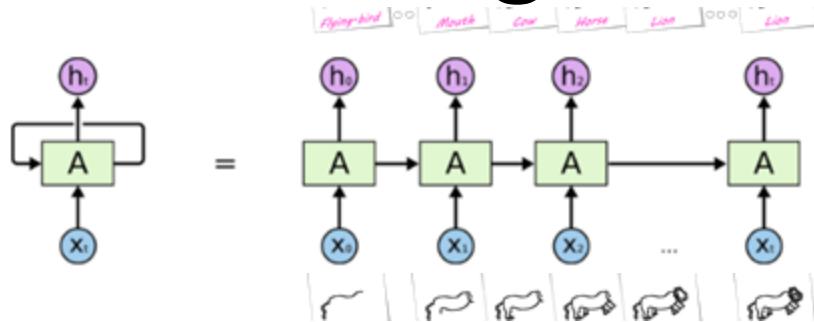


Recurrence “loop”

Error derivatives w.r.t weights in kth layer = $f(\text{ Error derivatives from (k+1)th layer })$
⇒ **Error derivatives in terms of themselves ??!!**



Modelling a RNN

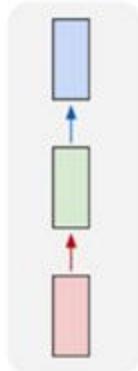


- RNN unrolled → Not “that different” from a CNN (feed-forward connections)
- But crucial differences exist !

Modelling a RNN

Recurrent Networks offer a lot of flexibility:

one to one

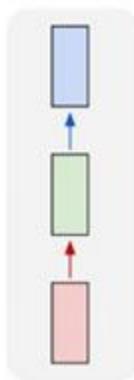


Vanilla Neural Networks

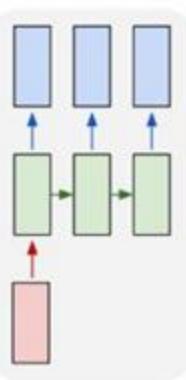
Modelling a RNN

Recurrent Networks offer a lot of flexibility:

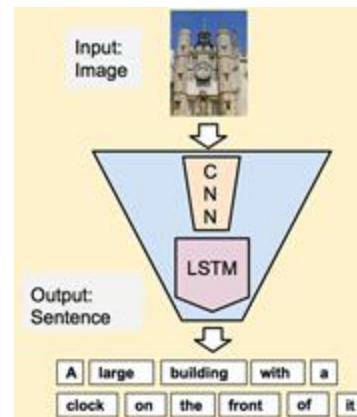
one to one



one to many



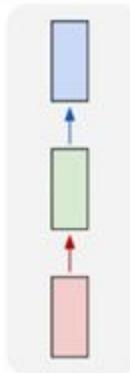
e.g. **Image Captioning**
image -> sequence of words



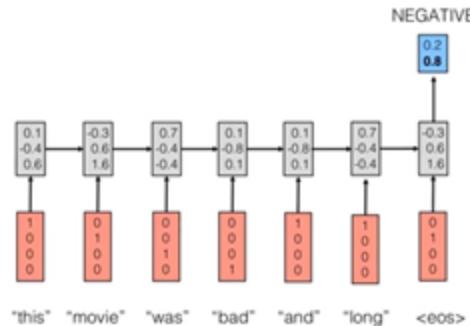
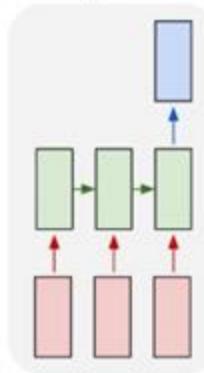
Modelling a RNN

Recurrent Networks offer a lot of flexibility:

one to one



many to one

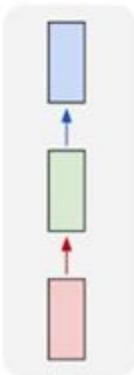


e.g. **Sentiment Classification**
sequence of words -> sentiment

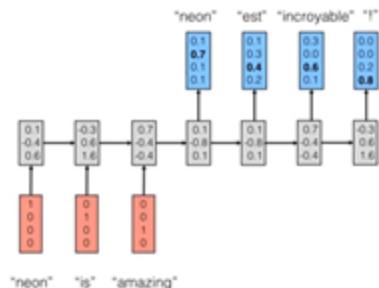
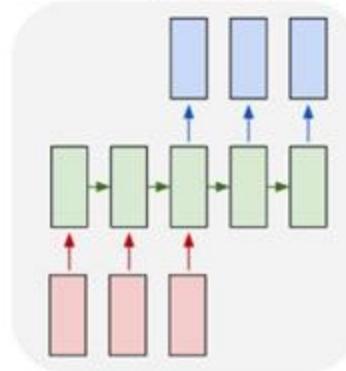
Modelling a RNN

Recurrent Networks offer a lot of flexibility:

one to one



many to many

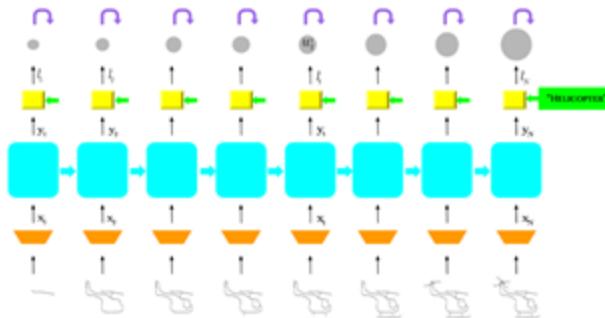
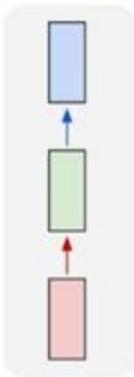


e.g. Machine Translation
seq of words -> seq of words

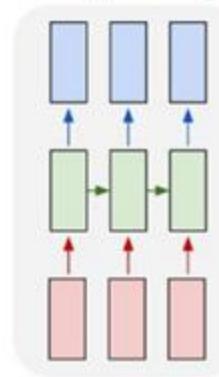
Modelling a RNN

Recurrent Networks offer a lot of flexibility:

one to one

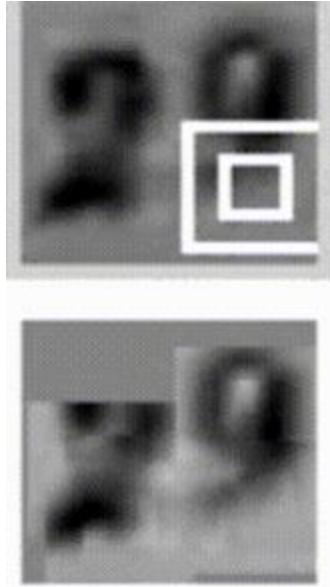


many to many



E.g. On-line freehand sketch recognition
Seq of strokes -> Seq of "guesses"

Paradigm: Sequential “processing”



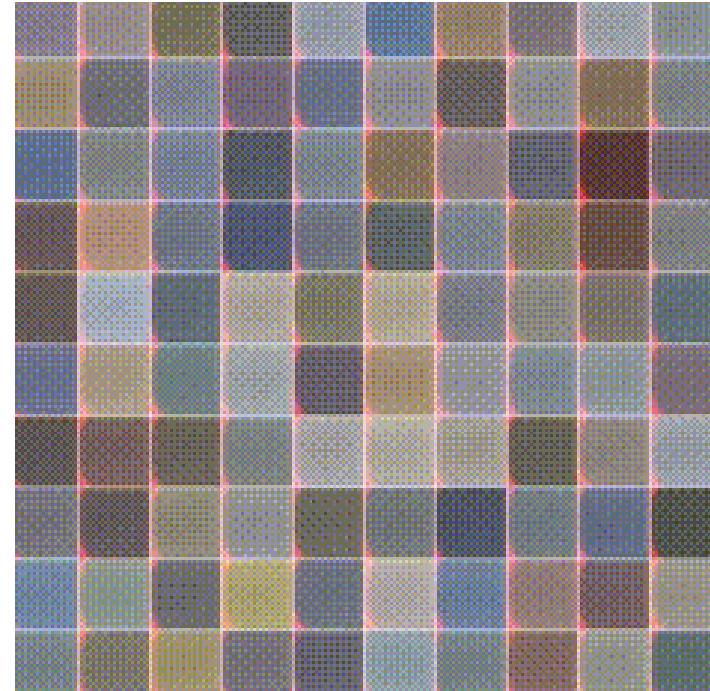
Reading house numbers

Input/Output may be
“fixed”, but
processing can be
sequential !

Multiple Object Recognition with Visual
Attention, Ba et al.

Paradigm: Sequential “processing”

DRAW: A Recurrent
Neural Network For
Image Generation,
Gregor et al.

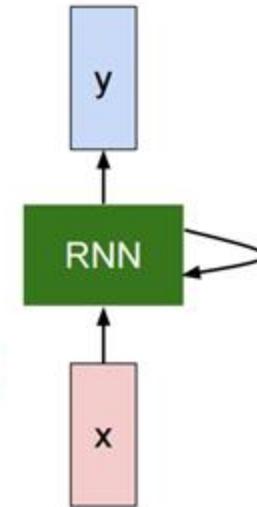


Recurrent Neural Network

We can process a sequence of vectors x by applying a recurrence formula at every time step:

$$h_t = f_W(h_{t-1}, x_t)$$

new state / old state input vector at some time step
some function with parameters W



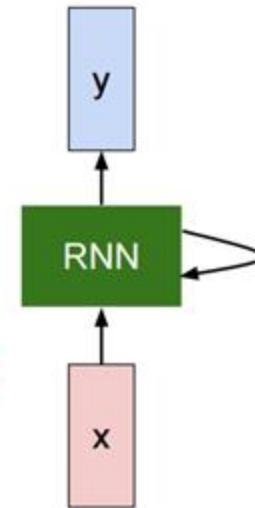
Task-specific abstraction for portion of sequence seen up to this point
 $(x_1, x_2 \dots x_{t-1})$

Recurrent Neural Network

We can process a sequence of vectors x by applying a recurrence formula at every time step:

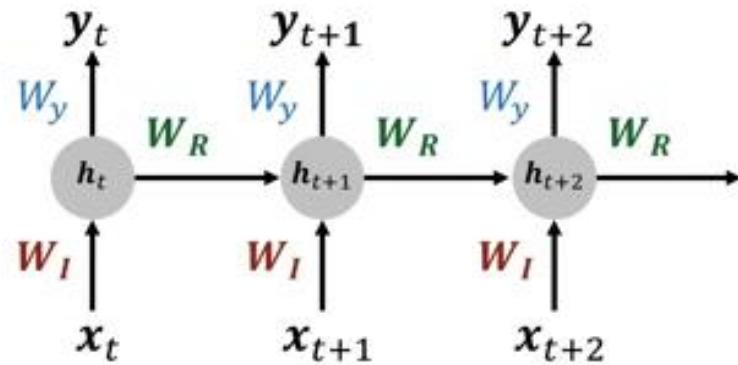
$$h_t = f_W(h_{t-1}, x_t)$$

new state / old state input vector at
 \ some time step
 some function
 with parameters W



Notice: the same function and the same set of parameters are used at every time step.

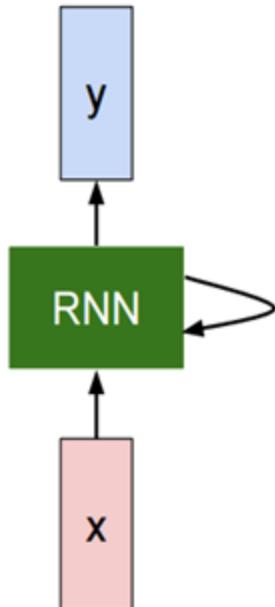
Recurrent Neural Network



3 sets of parameters - W_I, W_y, W_R (shared for each time-step)

(Vanilla) Recurrent Neural Network

The state consists of a single “*hidden*” vector \mathbf{h} :



$$h_t = f_W(h_{t-1}, x_t)$$



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

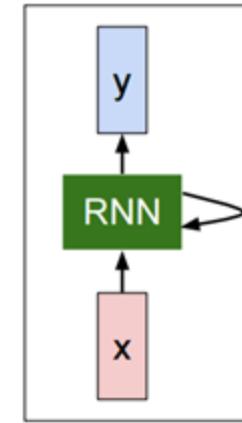
$$y_t = W_{hy}h_t$$

Recurrent Neural Network

**Character-level
language model
example**

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

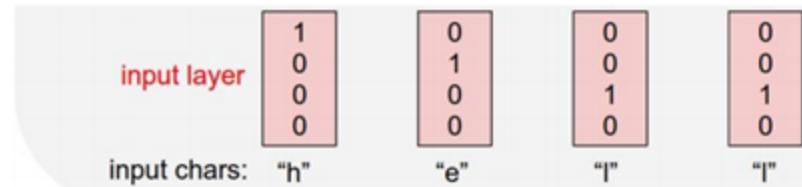


Recurrent Neural Network

**Character-level
language model
example**

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”



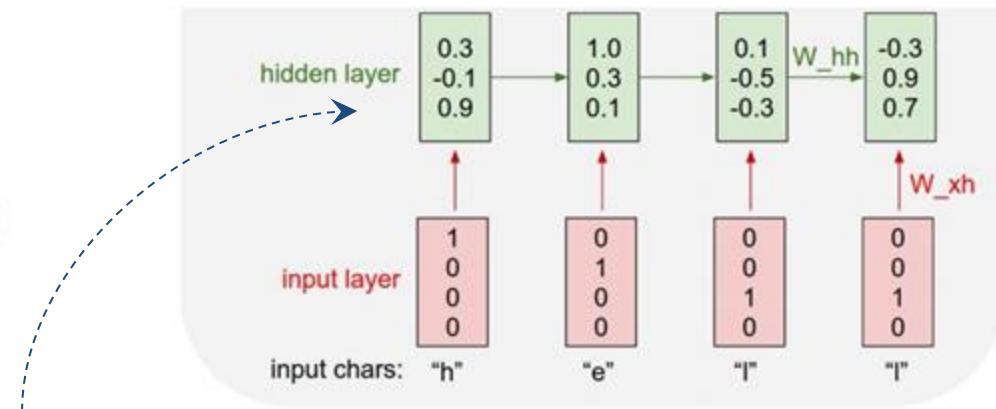
Recurrent Neural Network

Character-level
language model
example

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”



$h_t \rightarrow$ abstraction of current input upto step t

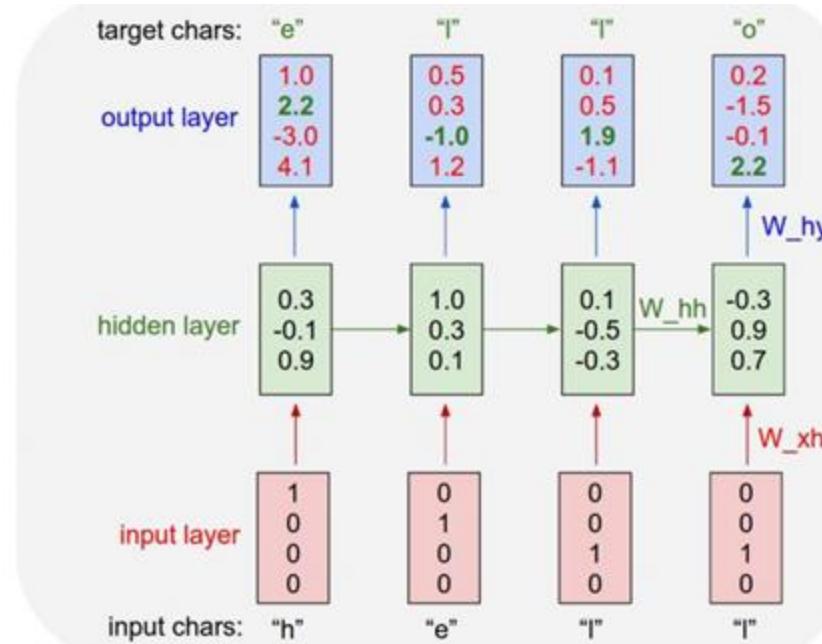
Recurrent Neural Network

$$y_t = W_{hy} h_t$$

**Character-level
language model
example**

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

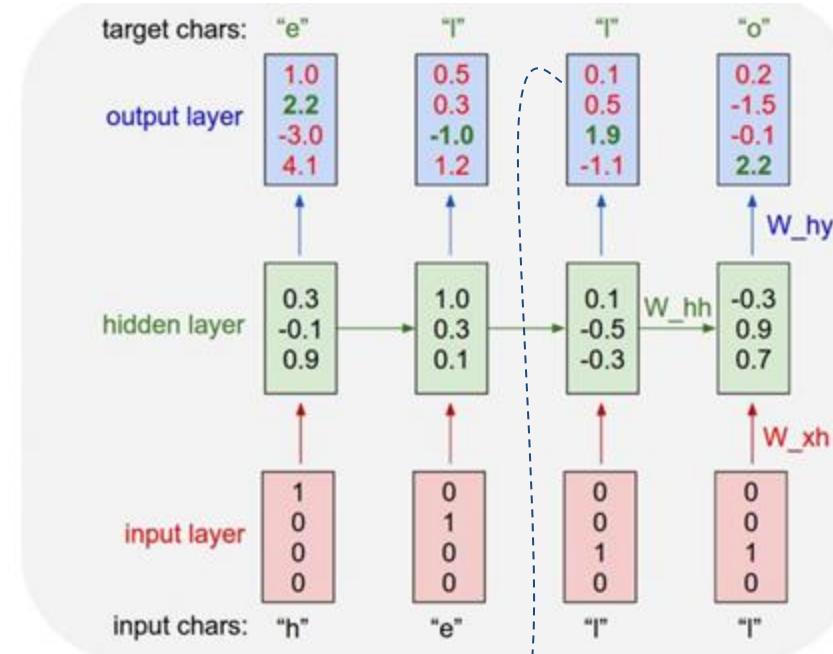


Recurrent Neural Network

**Character-level
language model
example**

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”



$$p(\text{next letter} = \text{"l"} | \text{previous letters} = \{\text{"h"}, \text{"e"}, \text{"l"}\})$$

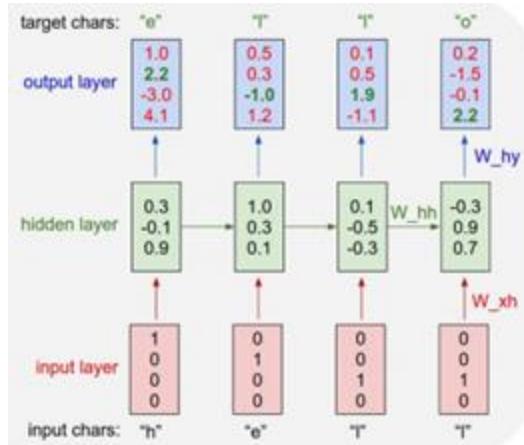
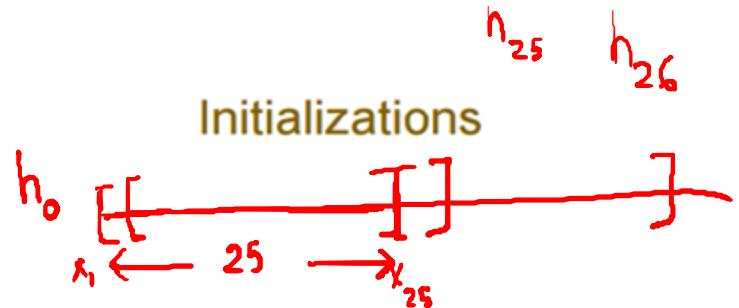
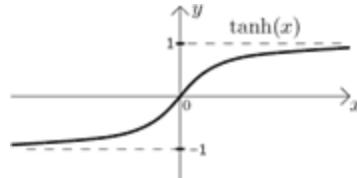
Recurrent Neural Network

```
# hyperparameters
hidden_size = 100 # size of hidden layer of neurons
seq_length = 25 = number of steps to unroll the RNN for
learning_rate = 1e-1

# model parameters
Wxh = np.random.randn(hidden_size, vocab_size)*0.01 # input to hidden
Whh = np.random.randn(hidden_size, hidden_size)*0.01 # hidden to hidden
Why = np.random.randn(vocab_size, hidden_size)*0.01 # hidden to output
bh = np.zeros((hidden_size, 1)) # hidden bias
by = np.zeros((vocab_size, 1)) # output bias
```

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$



Recurrent Neural Network

Main loop

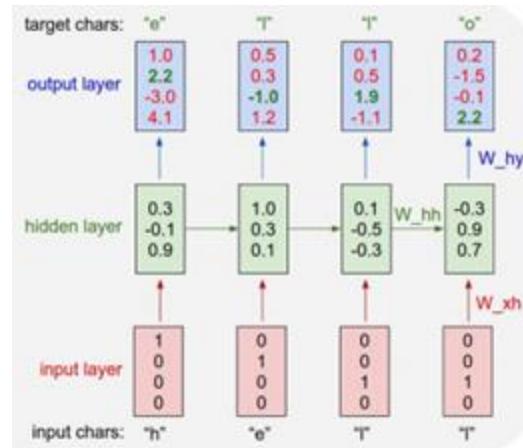
```
n, p = 0, 0
mwxh, mwhh, mwhy = np.zeros_like(Wxh), np.zeros_like(Whh), np.zeros_like(Why)
mbh, mbv = np.zeros_like(bh), np.zeros_like(by) # memory variables for Adagrad
smooth_loss = -np.log(1.0/vocab_size)*seq_length # loss at iteration 0
while True:
    # prepare inputs (we're sweeping from left to right in steps seq_length long)
    if p+seq_length+1 >= len(data) or n == 0:
        hprev = np.zeros((hidden_size,)) # reset RNN memory
        p = 0 # go from start of data
    inputs = [char_to_ix[ch] for ch in data[p:p+seq_length]] # input sequence
    targets = [char_to_ix[ch] for ch in data[p+1:p+seq_length+1]] # target sequence

    # sample from the model now and then
    if n % 100 == 0:
        sample_ix = sample(hprev, inputs[0], 200)
        txt = ''.join(ix_to_char[ix] for ix in sample_ix)
        print '----\n%s\n----' % (txt,)

    # forward seq_length characters through the net and fetch gradient
    loss, dxwh, dwhh, dwhy, dbh, dbv, hprev = lossFun(inputs, targets, hprev)
    smooth_loss = smooth_loss * 0.999 + loss * 0.001
    if n % 100 == 0: print 'iter %d, loss: %f' % (n, smooth_loss) # print progress

    # perform parameter update with Adagrad
    for param, dparam, mem in zip([Wxh, Whh, Why, bh, bv],
                                  [dxwh, dwhh, dwhy, dbh, dbv],
                                  [mwxh, mwhh, mwhy, mbh, mbv]):
        mem += dparam * dparam
        param += -learning_rate * dparam / np.sqrt(mem + 1e-8) # adagrad update

    p += seq_length # move data pointer
    n += 1 # iteration counter
```



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

Recurrent Neural Network

Main loop

```
n, p = 0, 0
mbch, mbhh, mbhy = np.zeros_like(wxh), np.zeros_like(whh), np.zeros_like(why)
mbh, mby = np.zeros_like(bh), np.zeros_like(by) # memory variables for Adagrad
smooth_loss = -np.log(1.0/vocab_size)*seq_length # loss at iteration 0
while True:
    # prepare inputs (we're sweeping from left to right in steps seq_length long)
    if p+seq_length+1 >= len(data) or n == 0:
        hprev = np.zeros((hidden_size,1)) # reset RNN memory
        p = 0 # go from start of data
    inputs = [char_to_ix[ch] for ch in data[p:p+seq_length]]
    targets = [char_to_ix[ch] for ch in data[p+1:p+seq_length+1]]

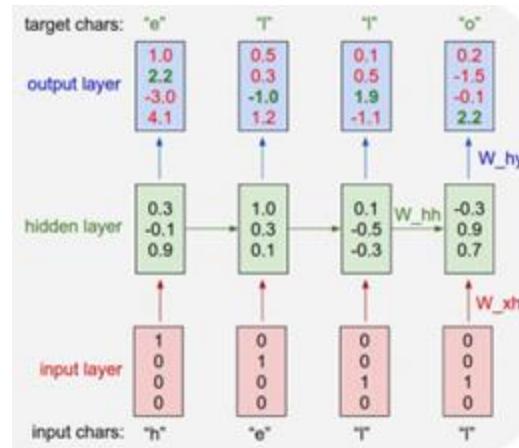
    # sample from the model now and then
    if n % 100 == 0:
        sample_ix = sample(hprev, inputs[0], 200)
        txt = ''.join(ix_to_char[ix] for ix in sample_ix)
        print '----\n%s\n----' % (txt,)

    # forward seq_length characters through the net and fetch gradient
    loss, dxh, dwhh, dwhy, dbh, dby, hprev = lossFun(inputs, targets, hprev)
    smooth_loss = smooth_loss * 0.999 + loss * 0.001
    if n % 100 == 0: print 'iter %d, loss: %f' % (n, smooth_loss) # print progress

    # perform parameter update with Adagrad
    for param, dparam, mem in zip([wxh, whh, why, bh, by],
                                  [dxh, dwhh, dwhy, dbh, dby],
                                  [mbch, mbhh, mbhy, mbh, mby]):
        mem += dparam * dparam
        param += -learning_rate * dparam / np.sqrt(mem + 1e-8) # adagrad update

    p += seq_length # move data pointer
    n += 1 # iteration counter
```

Forward prop



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

Recurrent Neural Network

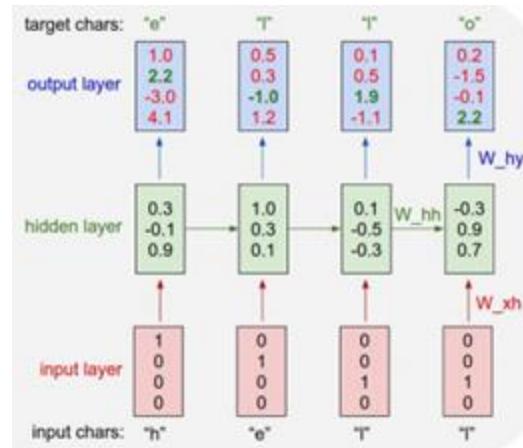
```
27 def lossFun(inputs, targets, hprev):
28     """
29     inputs,targets are both list of integers.
30     hprev is Hx1 array of initial hidden state
31     returns the loss, gradients on model parameters, and last hidden state
32     """
33     xs, hs, ys, ps = {}, {}, {}, {}
34     hs[-1] = np.copy(hprev)
35     loss = 0
36     # forward pass
37     for t in xrange(len(inputs)):
38         xs[t] = np.zeros((vocab_size,1)) # encode in 1-of-k representation
39         xs[t][inputs[t]] = 1
40         hs[t] = np.tanh(np.dot(Wxh, xs[t]) + np.dot(Whh, hs[t-1]) + bh) # hidden state
41         ys[t] = np.dot(Why, hs[t]) * by # unnormalized log probabilities for next chars
42         ps[t] = np.exp(ys[t]) / np.sum(np.exp(ys[t])) # probabilities for next chars
43         loss += -np.log(ps[t][targets[t],0]) # softmax (cross-entropy loss)
```

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

Softmax classifier

Forward prop



Recurrent Neural Network

Main loop

```

n, p = 0, 0
mwxh, mwhh, mwhy = np.zeros_like(wxh), np.zeros_like(whh), np.zeros_like(why)
mbh, mbv = np.zeros_like(bh), np.zeros_like(by) # memory variables for Adagrad
smooth_loss = -np.log(1.0/vocab_size)*seq_length # loss at iteration 0
while True:
    # prepare inputs (we're sweeping from left to right in steps seq_length long)
    if p+seq_length+1 >= len(data) or n == 0:
        hprev = np.zeros((hidden_size,)) # reset RNN memory
        p = 0 # go from start of data
    inputs = [char_to_ix[ch] for ch in data[p:p+seq_length]]
    targets = [char_to_ix[ch] for ch in data[p+1:p+seq_length+1]]

    # sample from the model now and then
    if n % 100 == 0:
        sample_ix = sample(hprev, inputs[0], 200)
        txt = ''.join(ix_to_char[ix] for ix in sample_ix)
        print '----\n' + txt + '----'

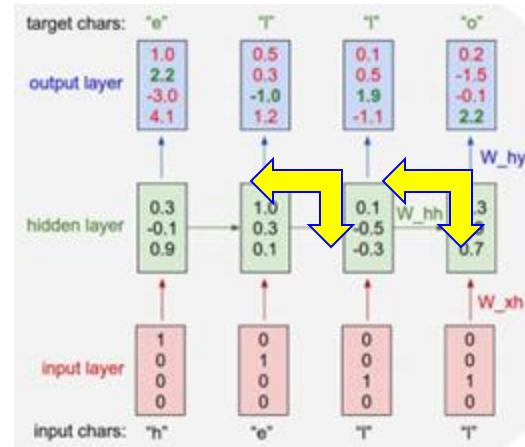
    # forward seq_length characters through the net and fetch gradient
    loss, dwxh, dwhh, dwhy, dbh, dbv, hprev = lossFun(inputs, targets, hprev)
    smooth_loss = smooth_loss * 0.999 + loss * 0.001
    if n % 100 == 0: print 'iter %d, loss: %f' % (n, smooth_loss) # print progress

    # perform parameter update with Adagrad
    for param, dparam, mem in zip([wxh, whh, why, bh, bv],
                                  [dwxh, dwhh, dwhy, dbh, dbv],
                                  [mwxh, mwhh, mwhy, mbh, mbv]):
        mem += dparam * dparam
        param += -learning_rate * dparam / np.sqrt(mem + 1e-8) # adagrad update

    p += seq_length # move data pointer
    n += 1 # iteration counter

```

Backprop



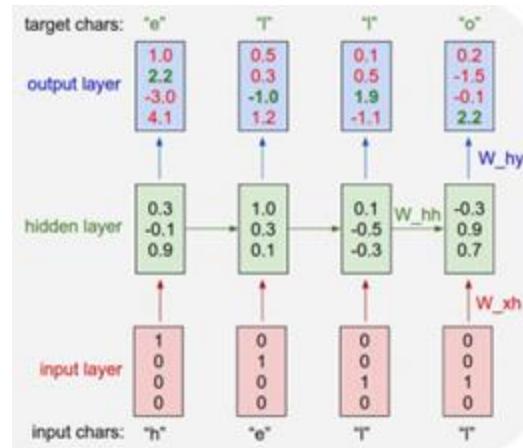
$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

Recurrent Neural Network

```
def sample(h, seed_ix, n):
    """
    sample a sequence of integers from the model
    h is memory state, seed_ix is seed letter for first time step
    """
    x = np.zeros((vocab_size, 1))
    x[seed_ix] = 1
    ixes = []
    for t in xrange(n):
        h = np.tanh(np.dot(Wxh, x) + np.dot(Whh, h) + bh)
        y = np.dot(Why, h) + by
        p = np.exp(y) / np.sum(np.exp(y))
        ix = np.random.choice(range(vocab_size), p=p.ravel())
        x = np.zeros((vocab_size, 1))
        x[ix] = 1
        ixes.append(ix)
    return ixes
```

Generation



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

Recurrent Neural Network

Sonnet 116 – Let me not ...

by William Shakespeare

Let me not to the marriage of true minds
Admit impediments. Love is not love
Which alters when it alteration finds,
Or bends with the remover to remove:
O no! it is an ever-fixed mark
That looks on tempests and is never shaken;
It is the star to every wandering bark,
Whose worth's unknown, although his height be taken.
Love's not Time's fool, though rosy lips and cheeks
Within his bending sickle's compass come:
Love alters not with his brief hours and weeks,
But bears it out even to the edge of doom.
If this be error and upon me proved,
I never writ, nor no man ever loved.

Recurrent Neural Network

at first:

tyntd-iafhatawiaoahrdemot lytdws e ,tfti, astai f ogoh eoase rrranbyne 'nhthnee e
plia tkldrgd t o idoe ns,smtt h ne etie h,hregtrs nigtike,aoaenns lng

↓ train more

"Tmont thithey" fomesscerliund
Keushey. Thom here
sheulke, anmerenith ol sivh I lalterthend Bleipile shuwyl fil on aseterlome
coaniogennc Phe lism thond hon at. MeiDimorotion in ther thize."

↓ train more

Aftair fall unsuch that the hall for Prince Velzonski's that me of
her hearly, and behs to so arwage fiving were to it beloge, pavu say falling misfort
how, and Gogition is so overelical and ofter.

↓ train more

"Why do what that day," replied Natasha, and wishing to himself the fact the
princess, Princess Mary was easier, fed in had oftened him.
Pierre aking his soul came to the packs and drove up his father-in-law women.

Recurrent Neural Network

open source textbook on algebraic geometry

The Stacks Project

home about tags explained tag lookup browse search bibliography recent comments blog add slogans

Browse chapters

Part	Chapter	online	TeX source	view pdf
Preliminaries	1. Introduction	online	tex	pdf
	2. Conventions	online	tex	pdf
	3. Set Theory	online	tex	pdf
	4. Categories	online	tex	pdf
	5. Topology	online	tex	pdf
	6. Sheaves on Spaces	online	tex	pdf
	7. Sites and Sheaves	online	tex	pdf
	8. Stacks	online	tex	pdf
	9. Fields	online	tex	pdf
	10. Commutative Algebra	online	tex	pdf

Parts

- [Preliminaries](#)
- [Schemes](#)
- [Topics in Scheme Theory](#)
- [Algebraic Spaces](#)
- [Topics in Geometry](#)
- [Deformation Theory](#)
- [Algebraic Stacks](#)
- [Miscellany](#)

Statistics

The Stacks project now consists of

- o 455910 lines of code
- o 14221 tags (56 inactive tags)
- o 2366 sections

Latex source

Recurrent Neural Network

Proof. Omitted. \square

Lemma 0.1. Let \mathcal{C} be a set of the construction.

Let \mathcal{C} be a gerber covering. Let \mathcal{F} be a quasi-coherent sheaves of \mathcal{O} -modules. We have to show that

$$\mathcal{O}_{\mathcal{O}_X} = \mathcal{O}_X(\mathcal{L})$$

Proof. This is an algebraic space with the composition of sheaves \mathcal{F} on X_{etale} we have

$$\mathcal{O}_X(\mathcal{F}) = \{morph_1 \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F})\}$$

where \mathcal{G} defines an isomorphism $\mathcal{F} \rightarrow \mathcal{G}$ of \mathcal{O} -modules. \square

Lemma 0.2. This is an integer \mathcal{Z} is injective.

Proof. See Spaces, Lemma ???. \square

Lemma 0.3. Let S be a scheme. Let X be a scheme and X is an affine open covering. Let $\mathcal{U} \subset X$ be a canonical and locally of finite type. Let X be a scheme. Let X be a scheme which is equal to the formal complex.

The following to the construction of the lemma follows.

Let X be a scheme. Let X be a scheme covering. Let

$$b : X \rightarrow Y' \rightarrow Y \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$

be a morphism of algebraic spaces over S and Y .

Proof. Let X be a nonzero scheme of X . Let X be an algebraic space. Let \mathcal{F} be a quasi-coherent sheaf of \mathcal{O}_X -modules. The following are equivalent

- (1) \mathcal{F} is an algebraic space over S .
- (2) If X is an affine open covering.

Consider a common structure on X and X the functor $\mathcal{O}_X(U)$ which is locally of finite type. \square

This since $\mathcal{F} \in \mathcal{F}$ and $x \in \mathcal{G}$ the diagram

$$\begin{array}{ccccc}
 S & \xrightarrow{\quad} & & & \\
 \downarrow & & & & \\
 \xi & \longrightarrow & \mathcal{O}_{X'} & \nearrow & \\
 & & \downarrow & & \\
 & & = \alpha' & \longrightarrow & X \\
 & & \downarrow & & \downarrow \\
 & & = \alpha' & \longrightarrow & \\
 & & & & Spec(K_\psi) \qquad \qquad \qquad Mor_{S_{etale}} \qquad d(\mathcal{O}_{X_{etale}}, \mathcal{G})
 \end{array}$$

is a limit. Then \mathcal{G} is a finite type and assume S is a flat and \mathcal{F} and \mathcal{G} is a finite type f_* . This is of finite type diagrams, and

- the composition of \mathcal{G} is a regular sequence,
- $\mathcal{O}_{X'}$ is a sheaf of rings.

Proof. We have see that $X = Spec(R)$ and \mathcal{F} is a finite type representable by algebraic space. The property \mathcal{F} is a finite morphism of algebraic stacks. Then the cohomology of X is an open neighbourhood of U . \square

Proof. This is clear that \mathcal{G} is a finite presentation, see Lemmas ???

A reduced above we conclude that U is an open covering of C . The functor \mathcal{F} is a "field"

$$\mathcal{O}_{X,x} \rightarrow \mathcal{F}_{X,x} : \mathcal{O}_{X_{etale}} \rightarrow \mathcal{O}_{X,x}^{-1} \mathcal{O}_{X,x}(\mathcal{O}_{X,x}^{\mathcal{F}})$$

is an isomorphism of covering of $\mathcal{O}_{X,x}$. If \mathcal{F} is the unique element of \mathcal{F} such that X is an isomorphism.

The property \mathcal{F} is a disjoint union of Proposition ?? and we can filtered set of presentations of a scheme \mathcal{O}_X -algebra with \mathcal{F} are opens of finite type over S . If \mathcal{F} is a scheme theoretic image points. \square

If \mathcal{F} is a finite direct sum $\mathcal{O}_{X,x}$ is a closed immersion, see Lemma ???. This is a sequence of \mathcal{F} is a similar morphism.

Recurrent Neural Network

The screenshot shows the GitHub repository page for `torvalds/linux`. The top navigation bar includes links for 'Explore', 'Gist', 'Blog', and 'Help'. The repository owner is `karpathy`. Key statistics displayed are 3,711 watches, 23,054 stars, and 9,141 forks. The repository name is `linux`, and the current branch is `master`. The main content area shows a list of recent commits, with the most recent being a merge from the `'drm-fixes'` branch of `git://people.freedesktop.org/~airlied/linux`, authored by `torvalds` 9 hours ago. Other commits listed include merges from `Documentation`, `arch`, `block`, `crypto`, `drivers`, `firmware`, `fs`, `include`, and `init` branches. On the right side, there are links for 'Code' (74 pull requests), 'Pulse', 'Graphs', and download options ('HTTPS clone URL', 'Clone In Desktop', 'Download ZIP').

Linux kernel source tree

520,037 commits 1 branch 420 releases 5,039 contributors

branch: master / +

Merge branch 'drm-fixes' of git://people.freedesktop.org/~airlied/linux ...

torvalds authored 9 hours ago latest commit 4b1706927d

Branch	Commit Message	Time Ago
Documentation	Merge git://git.kernel.org/pub/scm/linux/kernel/git/torvalds/... pending	6 days ago
arch	Merge branch 'x86-urgent-for-linus' of git://git.kernel.org/pub/scm/...	a day ago
block	block: discard bdi_unregister() in favour of bdi_destroy()	9 days ago
crypto	Merge git://git.kernel.org/pub/scm/linux/kernel/git/herbert/crypto-2.6	10 days ago
drivers	Merge branch 'drm-fixes' of git://people.freedesktop.org/~airlied/linux	9 hours ago
firmware	firmware/hex2hw.c: restore missing default in switch statement	2 months ago
fs	vfs: read file_handle only once in handle_to_path	4 days ago
include	Merge branch 'perf-urgent-for-linus' of git://git.kernel.org/pub/scm/...	a day ago
init	init: fix regression by supporting devices with major:minor:offset fo...	a month ago
scripts	scripts: remove 'grep -q' and add -Egrep because -q does not do what we want	1 year ago

Pull requests: 74

Pulse

Graphs

HTTPS clone URL: <https://github.com/torvalds/linux>

You can clone with HTTPS, SSH, or Subversion.

Clone In Desktop

Download ZIP

Recurrent Neural Network

```
static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << i))
            pipe = (in_use & UMXTHREAD_UNCCA) +
                ((count & 0x00000000fffff8) & 0x000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &soffset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

Generated
C code



Recurrent Neural Network

```
/*
 * Copyright (c) 2006-2010, Intel Mobile Communications. All rights reserved.
 *
 * This program is free software; you can redistribute it and/or modify it
 * under the terms of the GNU General Public License version 2 as published by
 * the Free Software Foundation.
 *
 *      This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
 *
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software Foundation,
 * Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 */

#include <linux/kexec.h>
#include <linux/errno.h>
#include <linux/io.h>
#include <linux/platform_device.h>
#include <linux/multi.h>
#include <linux/ckevent.h>

#include <asm/io.h>
#include <asm/prom.h>
#include <asm/e820.h>
#include <asm/system_info.h>
#include <asm/seteov.h>
#include <asm/pgproto.h>
```

Recurrent Neural Network

Main loop

```
n, p = 0, 0
mwxh, mwhh, mwyh = np.zeros_like(wxh), np.zeros_like(whh), np.zeros_like(why)
mbh, mbw = np.zeros_like(bh), np.zeros_like(by) # memory variables for Adagrad
smooth_loss = -np.log(1.0/vocab_size)*seq_length # loss at iteration 0
while True:
    # prepare inputs (we're sweeping from left to right in steps seq_length long)
    if p+seq_length+1 >= len(data) or n == 8:
        hprev = np.zeros((hidden_size, 1)) # reset RNN memory
        p = 0 # go from start of data
    inputs = [char_to_ix[ch] for ch in data[p:p+seq_length]]
    targets = [char_to_ix[ch] for ch in data[p+1:p+seq_length+1]]

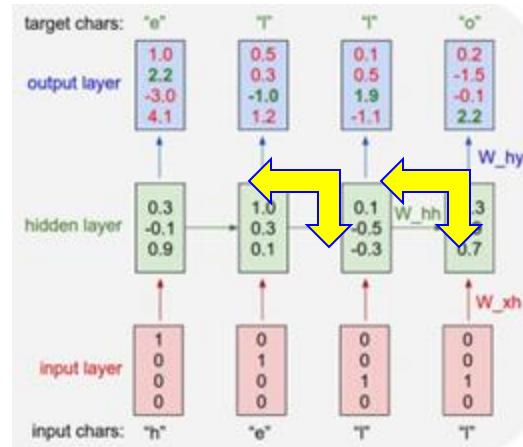
    # sample from the model now and then
    if n % 100 == 0:
        sample_ix = sample(hprev, inputs[0], 200)
        txt = ''.join(ix_to_char[ix] for ix in sample_ix)
        print '----\n% %----' % (txt, )

    # forward seq_length characters through the net and fetch gradient
    loss, dwxh, dwhh, dwhy, dbh, dbw, hprev = lossFun(inputs, targets, hprev)
    smooth_loss = smooth_loss * 0.999 + loss * 0.001
    if n % 100 == 0: print 'iter %d, loss: %f' % (n, smooth_loss) # print progress

    # perform parameter update with Adagrad
    for param, dparam, mem in zip([wxh, whh, why, bh, bw],
                                  [dwxh, dwhh, dwhy, dbh, dbw],
                                  [mwxh, mwhh, mwyh, mbh, mbw]):
        mem += dparam * dparam
        param += -learning_rate * dparam / np.sqrt(mem + 1e-8) # adagrad update

    p += seq_length # move data pointer
    n += 1 # iteration counter
```

Backprop

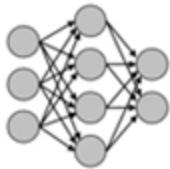


$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

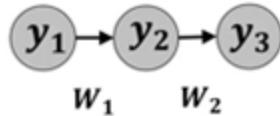
Training RNNs

- Recap: Training CNN



$$y_i = g \left(\sum_j W_{ij} x_j + b_i \right)$$

$$\frac{\partial C}{\partial W} = \frac{\partial C}{\partial g} \cdot \frac{\partial g}{\partial a} \cdot \frac{\partial a}{\partial W}$$



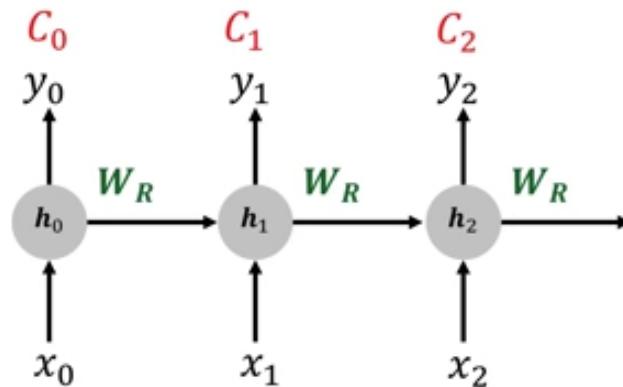
$$y_k = g(\underbrace{W y_{k-1} + b}_{\text{Recurrence}})$$

- Single, final cost C
- Non-shared weights (W_1, W_2)

Training RNNs

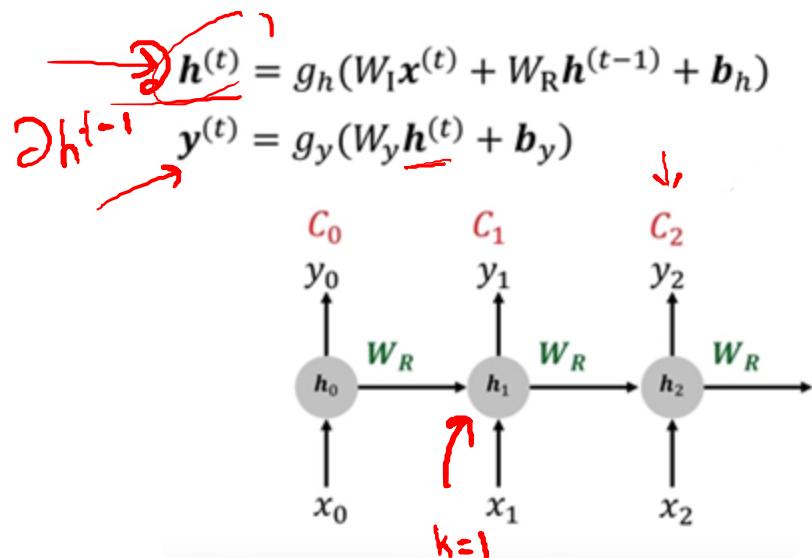
$$\mathbf{h}^{(t)} = g_h(W_{\text{I}} \mathbf{x}^{(t)} + W_{\text{R}} \mathbf{h}^{(t-1)} + \mathbf{b}_h)$$

$$\mathbf{y}^{(t)} = g_y(W_y \mathbf{h}^{(t)} + \mathbf{b}_y)$$



- (Potentially) multiple, intermediate costs C_0, C_1, C_2
- Shared weights W_R

Training RNNs - Backpropagation through time



$$\frac{\partial C}{\partial W_R} = \sum_{t=1}^T \frac{\partial C_t}{\partial W_R}$$

$$\frac{\partial C_t}{\partial W_R} = \sum_{k=1}^t \frac{\partial C_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W_R}$$

$$\frac{\partial h_t}{\partial h_k} = \prod_{i=k+1}^t \frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=k+1}^t W_R^T \text{diag}(g'_h(W_I x^{(i)} + W_R h^{(i-1)} + b_h))$$

$$\frac{\partial C_t}{\partial W_R} = \frac{\partial C_t}{\partial h_k} \cdot \frac{\partial h_t}{\partial h_{k-1}} \cdot \frac{\partial h_{k-1}}{\partial h_{k-2}} \dots$$

- (Potentially) multiple, intermediate costs C_0, C_1, C_2
- Shared weights W_R

Training RNNs - BPTT, Vanishing gradients

$$\mathbf{h}^{(t)} = g_h(W_I \mathbf{x}^{(t)} + W_R \mathbf{h}^{(t-1)} + \mathbf{b}_h)$$

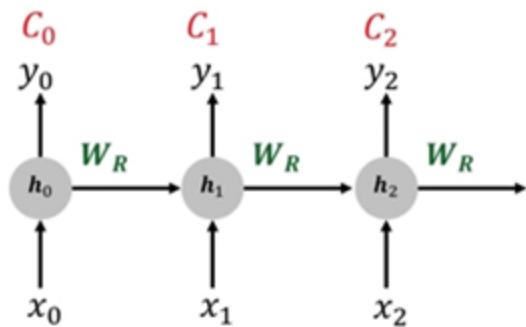
$$\mathbf{y}^{(t)} = g_y(W_y \mathbf{h}^{(t)} + \mathbf{b}_y)$$

$$\frac{\partial C}{\partial W_R} = \sum_{t=1}^T \frac{\partial C_t}{\partial W_R}$$

$$\frac{\partial C_t}{\partial W_R} = \sum_{k=1}^t \frac{\partial C_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W_R}$$

$$\frac{\partial h_t}{\partial h_k} = \prod_{i=k+1}^t \frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=k+1}^t W_R^T \text{diag}(g'_h(W_I x^{(i)} + W_R h^{(i-1)} + b_h))$$

$$\left\| \frac{\partial h_i}{\partial h_{i-1}} \right\| \leq \|W_R^T\| \|\text{diag}(g'_h(W_I x^{(i)} + W_R h^{(i-1)} + b_h))\| \leq \gamma_{W_R} \gamma_{g_h}$$



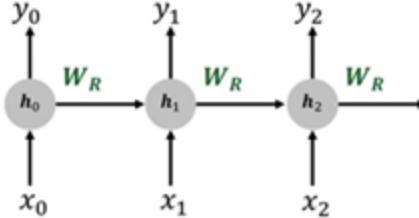
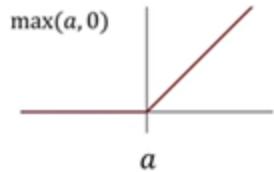
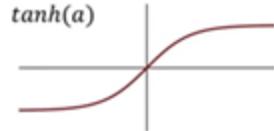
Vanishing gradients !

(Pascanu, et al., *On the difficulty of training Recurrent Neural Networks.*)

Training RNNs

Vanishing/Exploding gradients

Choice of $a(.)$
is crucial too!



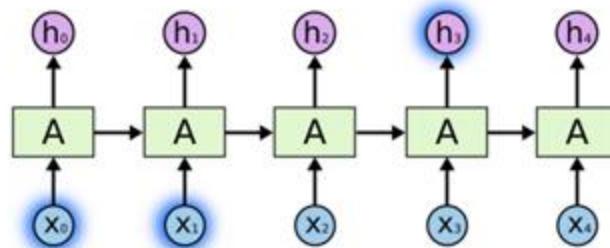
$$\frac{\partial C_{100}}{\partial W_R} = \frac{\partial C_{100}}{\partial y_{100}} \dots W_R \frac{\partial g_{100}}{\partial a_{100}} \dots W_R \frac{\partial g_{99}}{\partial a_{99}} \dots$$

$$\frac{\partial C_T}{\partial W_R} \propto |W_R|^T \left| \frac{\partial g}{\partial a} \right|^T$$

- $|W_R| \ll 1 \rightarrow$ Vanishing gradients
- $|W_R| \gg 1 \rightarrow$ Exploding gradients
- $|W_R|$ 'close to 1' is ideal

Fundamental issues with RNNs

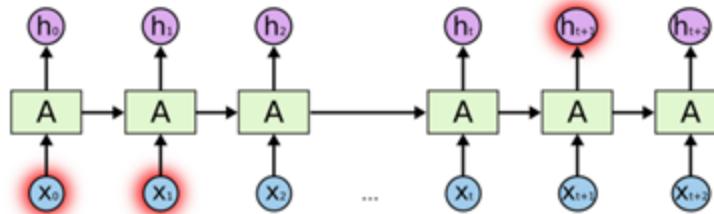
- RNNs connect previous information to the present task
 - Previous video frames may help understand present frame
- Sometimes we need only look at recent previous information to predict
 - To predict the last word of “The clouds are in the *sky*” we don’t need any further context. It is obvious that the word is “*sky*”

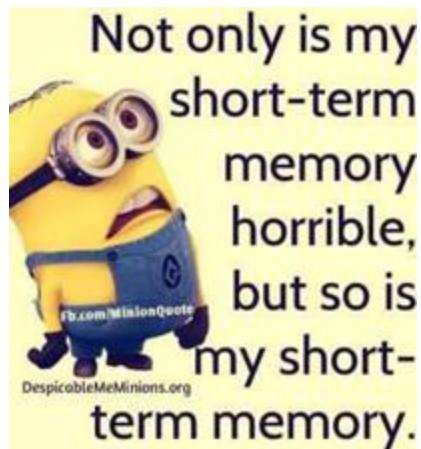


Fundamental issues with RNNs

Problem of Long-term dependency

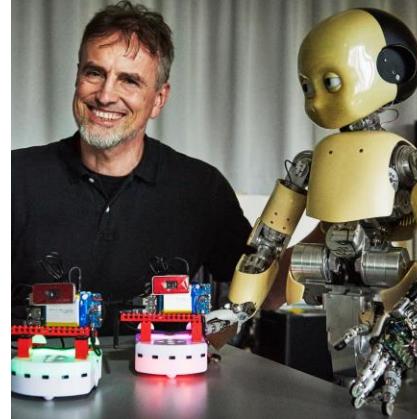
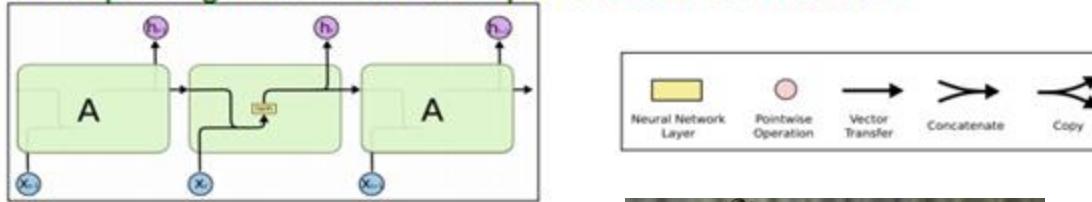
- There are cases where we need more context
 - To predict the last word in the sentence “I grew up in France...I speak fluent *French*”
 - Using only recent information suggests that the last word is the name of a language. But more distant past indicates that it is French
 - It is possible that the gap between the relevant information and where it is needed is very large





Long Short Term Memory (LSTM)

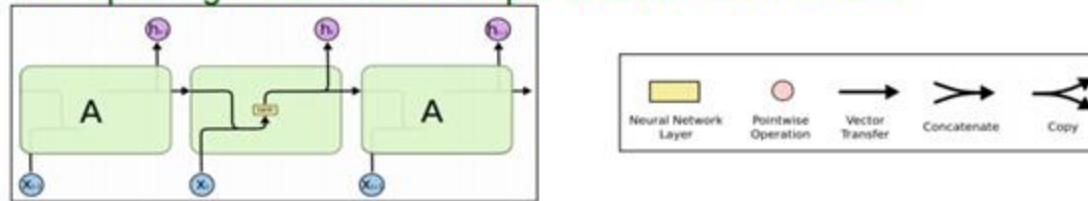
- Explicitly designed to avoid the long-term dependency problem
- RNNs have the form of a repeating chain structure
 - The repeating module has a simple structure such as tanh



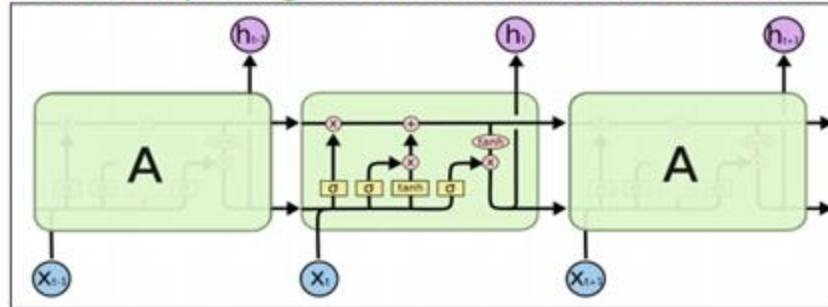
Proposed by Juergen Schmidhuber (1997 !)

Long Short Term Memory (LSTM)

- Explicitly designed to avoid the long-term dependency problem
- RNNs have the form of a repeating chain structure
 - The repeating module has a simple structure such as tanh



- LSTMs also have a chain structure
 - but the repeating module has a different structure

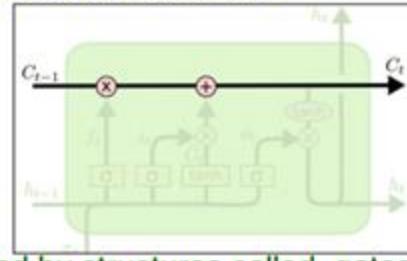


LSTM

Core idea behind LSTM

- The key to LSTM is the cell state, C_t , the horizontal line running through the top of the diagram
- Like a conveyor belt

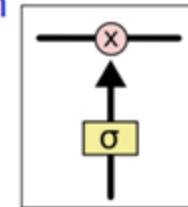
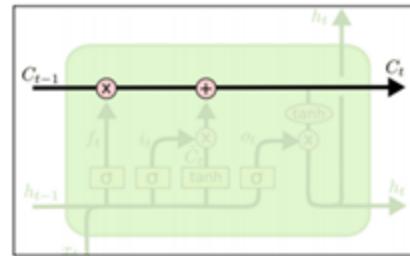
- Runs through entire chain with minor interactions
- LSTM does have the ability to remove/add information to cell state regulated by structures called gates



LSTM

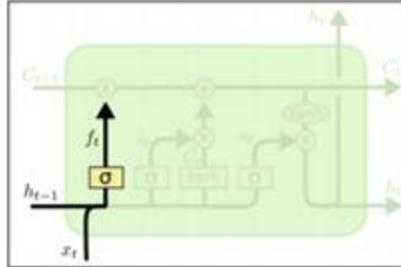
Core idea behind LSTM

- The key to LSTM is the cell state, C_t , the horizontal line running through the top of the diagram
- Like a conveyor belt
 - Runs through entire chain with minor interactions
 - LSTM does have the ability to remove/add information to cell state regulated by structures called gates
- Gates are an optional way to let information through
- Consist of a sigmoid and a multiplication operation
- Sigmoid outputs a value between 0 and 1
 - 0 means let nothing through
 - 1 means let everything through
- LSTM has three of these gates, to protect and control cell state



LSTM - Components

- Example of language model: predict next word based on previous ones
 - Cell state may include the gender of the present subject
- First step: information to throw away from cell state



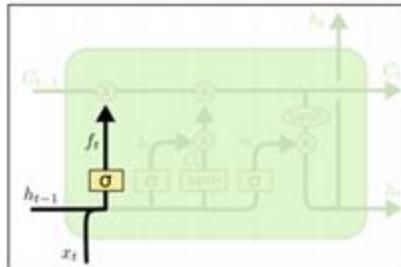
Called *forget gate layer*

It looks at h_{t-1} and x_t and outputs a number between 0 and 1 for each member of C_{t-1} for whether to forget

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

LSTM - Components

- Example of language model: predict next word based on previous ones
 - Cell state may include the gender of the present subject
- First step: information to throw away from cell state



Called *forget gate layer*

It looks at h_{t-1} and x_t and outputs a number between 0 and 1 for each member of C_{t-1} for whether to forget

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

In language model

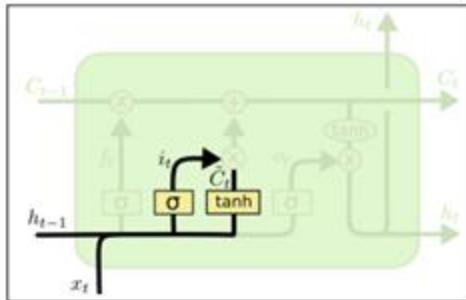
consider trying to predict the next word based on all previous ones

The cell state may include the gender of the present subject
so that the proper pronouns can be used

When we see a new subject we want to forget old subject

LSTM - Components

- Next step is to decide as to what new information we're going to store in the cell state



This has two parts:

first a sigmoid layer called *Input gate layer*: decides which values we will update
Next a tanh layer creates a vector of new candidate values \tilde{C}_t that could be added to the state.

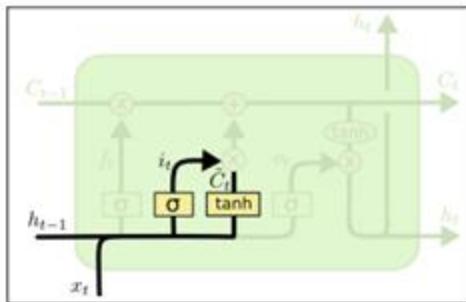
In the third step we will combine these two to create an update to the state

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

LSTM - Components

- Next step is to decide as to what new information we're going to store in the cell state



This has two parts:

first a sigmoid layer called *Input gate layer*: decides which values we will update
Next a tanh layer creates a vector of new candidate values \tilde{C}_t that could be added to the state.

In the third step we will combine these two to create an update to the state

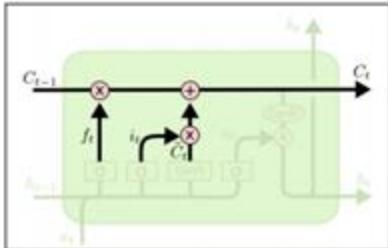
In the Language model, we'd want to add the gender of the new subject to the cell state, to replace the old One we are forgetting

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

LSTM - Components

- It's now time to update old cell state C_{t-1} into new cell state C_t
 - The previous step decided what we need to do
 - We just need to do it



We multiply the old state by f_t , forgetting the things we decided to forget earlier.
Then we add $i_t * \tilde{C}_t$
This is the new candidate values, scaled by how much we decided to update each state value

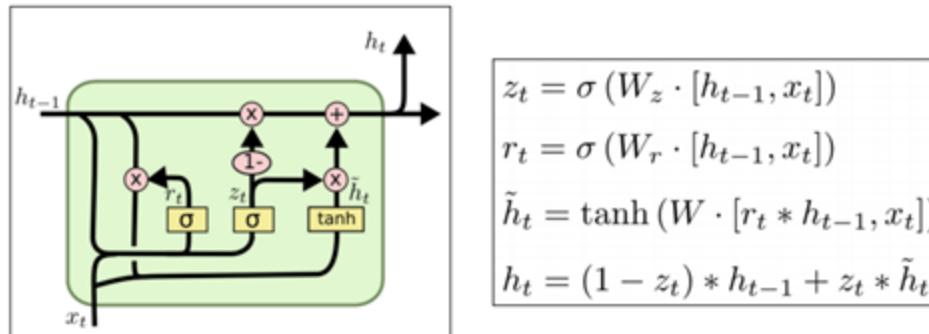
In the Language model, this is where we'd actually drop the information about the old subject's gender and add the new information, as we decided in previous steps

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

GRU: a LSTM variant

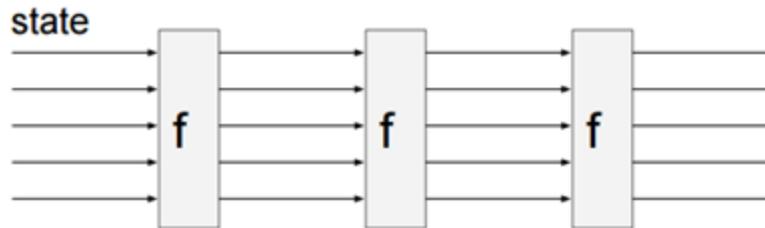
Gated Recurrent Unit (GRU)

- A dramatic variant of LSTM
 - It combines the forget and input gates into a single update gate
 - It also merges the cell state and hidden state, and makes some other changes
 - The resulting model is simpler than LSTM models
 - Has become increasingly popular

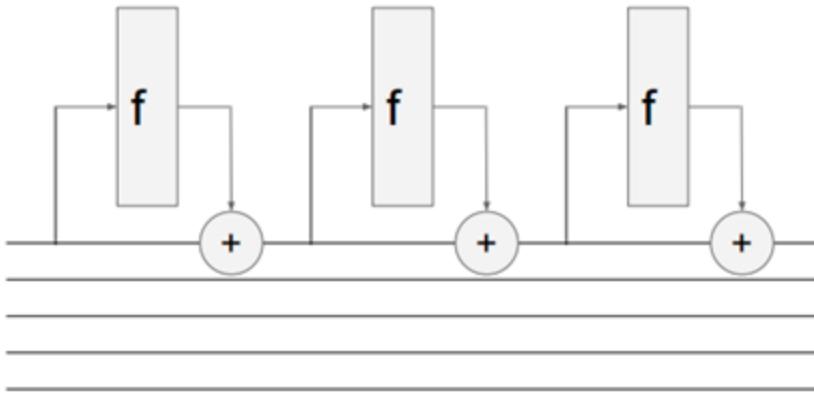


Long Short Term Memory (LSTM)

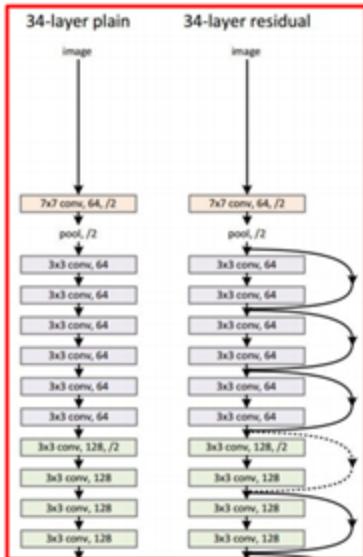
RNN



LSTM
(ignoring
forget gates)

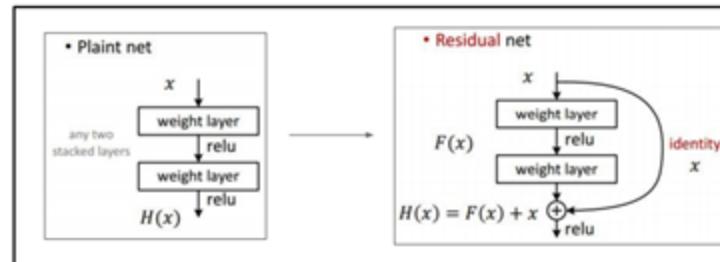


Long Short Term Memory (LSTM)



Recall:
“PlainNets” vs. ResNets

ResNet is to PlainNet what LSTM is to RNN, kind of.

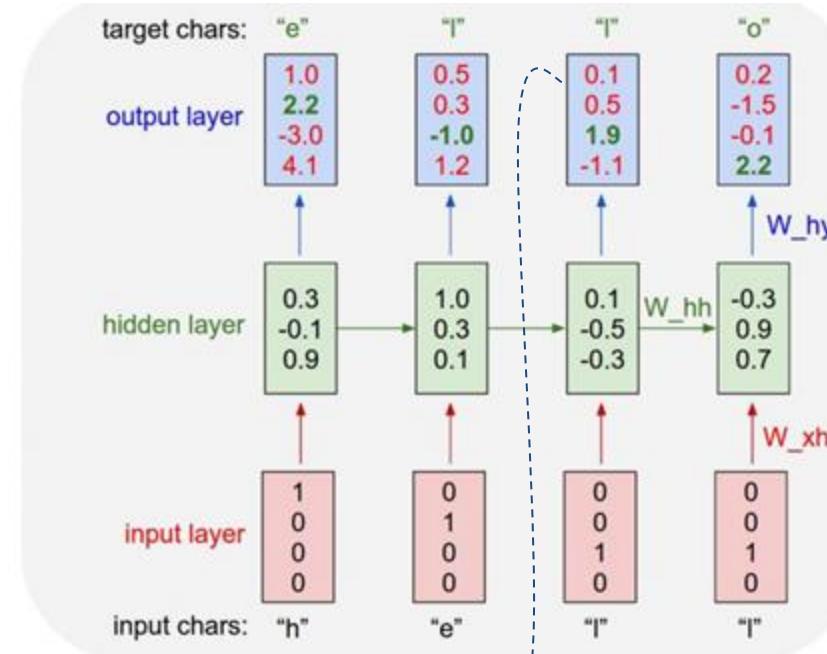


Recurrent Neural Network

**Character-level
language model
example**

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”



$$p(\text{next letter} = \text{"l"} | \text{previous letters} = \{\text{"h"}, \text{"e"}, \text{"l"}\})$$

Recurrent Neural Network

Searching for interpretable cells

"You mean to imply that I have nothing to eat out of.... On the contrary, I can supply you with everything even if you want to give dinner parties," warmly replied Chichagov, who tried by every word he spoke to prove his own rectitude and therefore imagined Kutuzov to be animated by the same desire.

Kutuzov, shrugging his shoulders, replied with his subtle penetrating smile: "I meant merely to say what I said."

quote detection cell

Recurrent Neural Network

Searching for interpretable cells

Cell sensitive to position in line:

```
The sole importance of the crossing of the Berezina lies in the fact  
that it plainly and indubitably proved the fallacy of all the plans for  
cutting off the enemy's retreat and the soundness of the only possible  
line of action--the one Kutuzov and the general mass of the army  
demanded--namely, simply to follow the enemy up. The French crowd fled  
at a continually increasing speed and all its energy was directed to  
reaching its goal. It fled like a wounded animal and it was impossible  
to block its path. This was shown not so much by the arrangements it  
made for crossing as by what took place at the bridges. When the bridges  
broke down, unarmed soldiers, people from Moscow and women with children  
who were with the French transport, all--carried on by vis inertiae--  
pressed forward into boats and into the ice-covered water and did not,  
surrender.
```

line length tracking cell

Recurrent Neural Network

Searching for interpretable cells

```
/* Duplicate LSM field information. The lsm_rule is opaque, so
 * re-initialized. */
static inline int audit_dupe_lsm_field(struct audit_field *df,
    struct audit_field *sf)
{
    int ret = 0;
    char *lsm_str;
    /* our own copy of lsm_str */
    lsm_str = kstrdup(sf->lsm_str, GFP_KERNEL);
    if (unlikely(!lsm_str))
        return -ENOMEM;
    df->lsm_str = lsm_str;
    /* our own (refreshed) copy of lsm_rule */
    ret = security_audit_rule_init(df->type, df->op, df->lsm_str,
        (void *) &df->lsm_rule);
    /* Keep currently invalid fields around in case they
     * become valid after a policy reload. */
    if (ret == -EINVAL) {
        pr_warn("audit rule for LSM \\'%s\\' is invalid\\n",
            df->lsm_str);
        ret = 0;
    }
    return ret;
}
```

quote/comment cell

Applications of RNNs in CV

- Image/Video Captioning
- Visual Question Answering
- Action / Activity Recognition from video
- See
<https://github.com/kjw0612/awesome-rnn> for a larger list ...

Applications of RNNs in CV

- **Image/Video Captioning**
- Visual Question Answering
- Action / Activity Recognition
from video

Image Captioning



man in black shirt is playing guitar.



construction worker in orange safety vest is working on road.



two young girls are playing with lego toy.

(Slides by Marc Bolaños): Karpathy, Andrej, and Li Fei-Fei. "Deep visual-semantic alignments for generating image descriptions." CVPR 2015

Recap: Language generation

Training this on a lot of sentences would give us a language model. A way to predict

$P(\text{next word} \mid \text{previous words})$

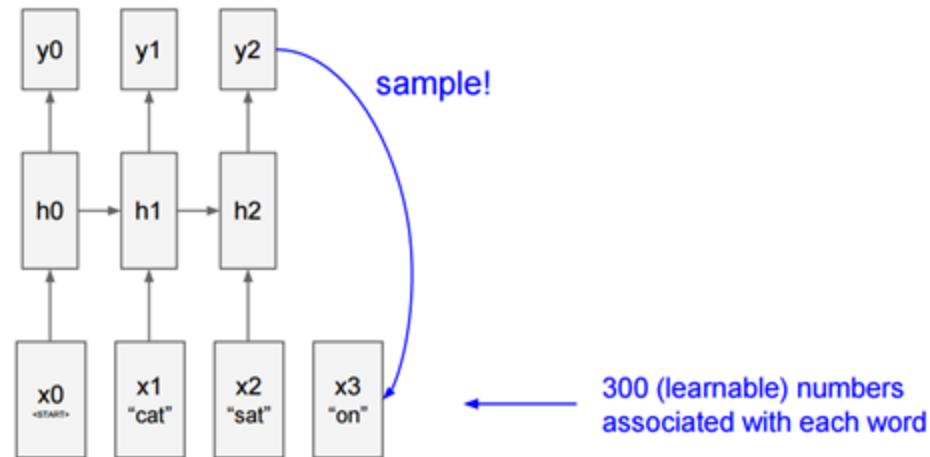
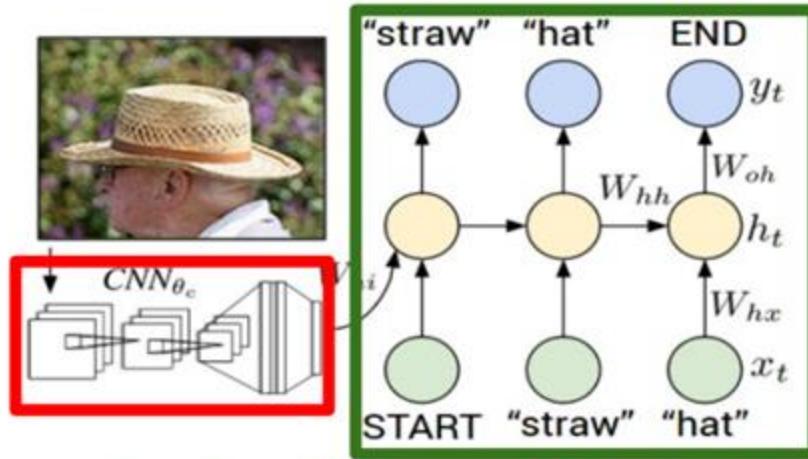


Image Captioning

Recurrent Neural Network



Convolutional Neural Network

Image Captioning

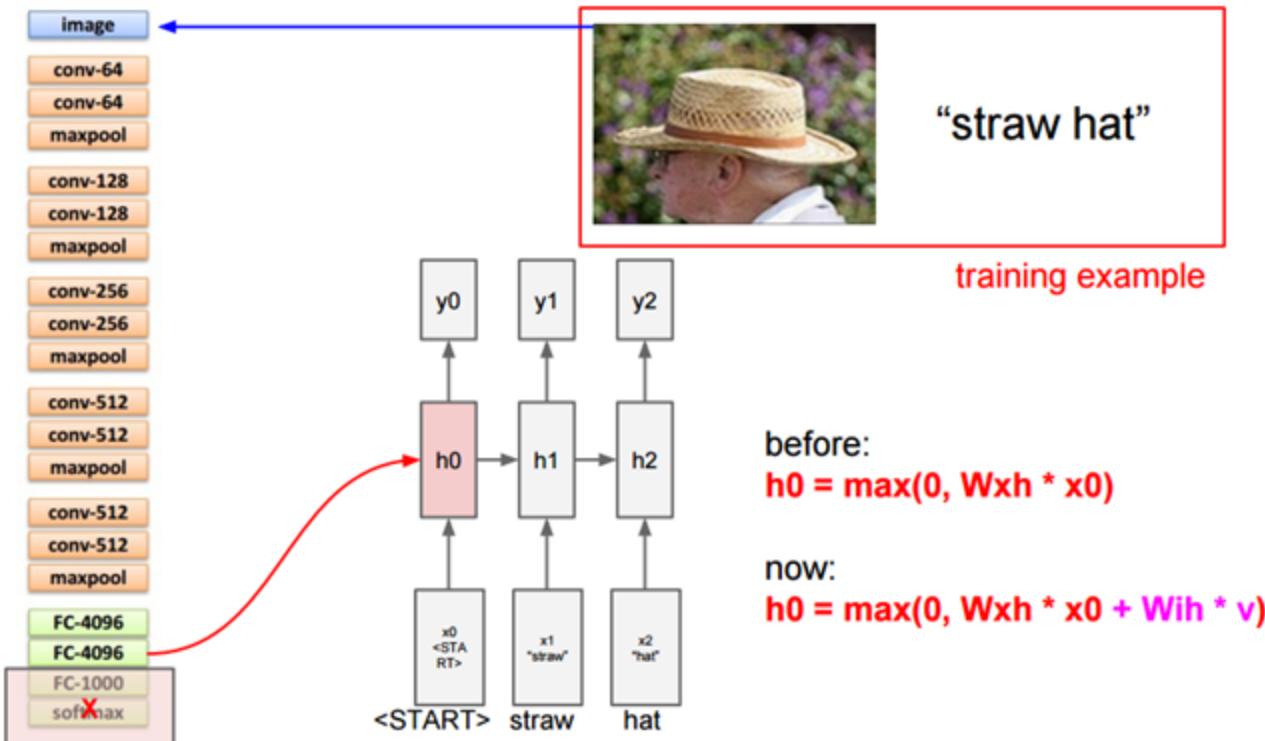


Image Captioning

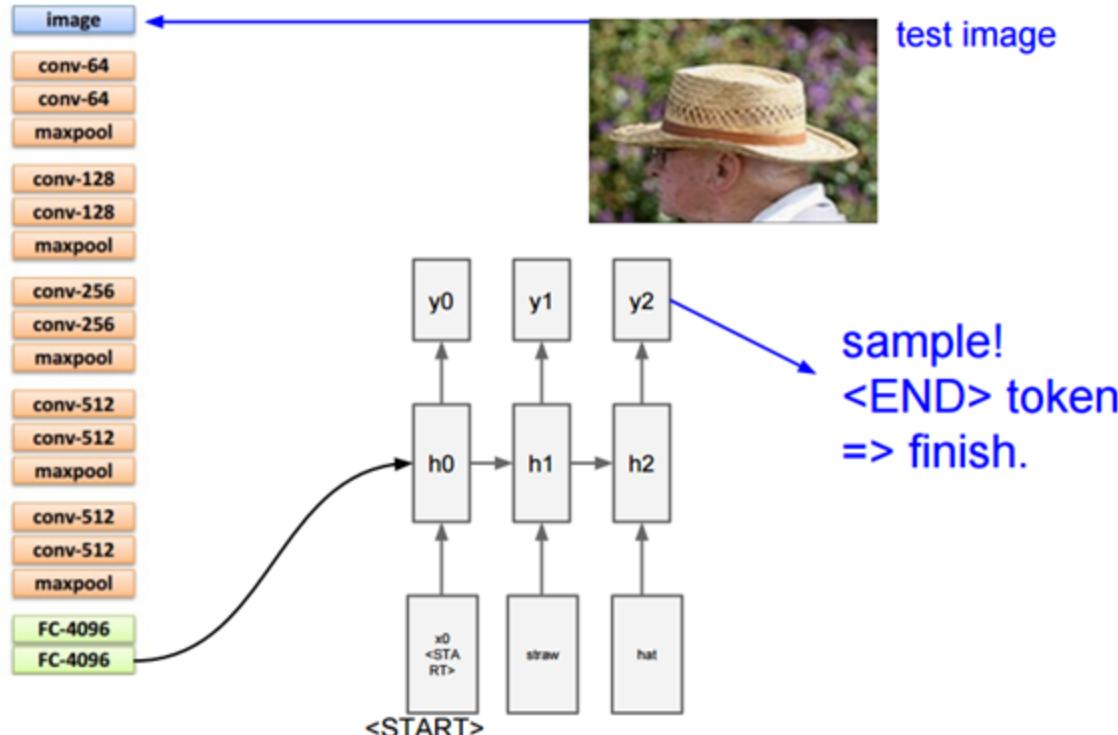
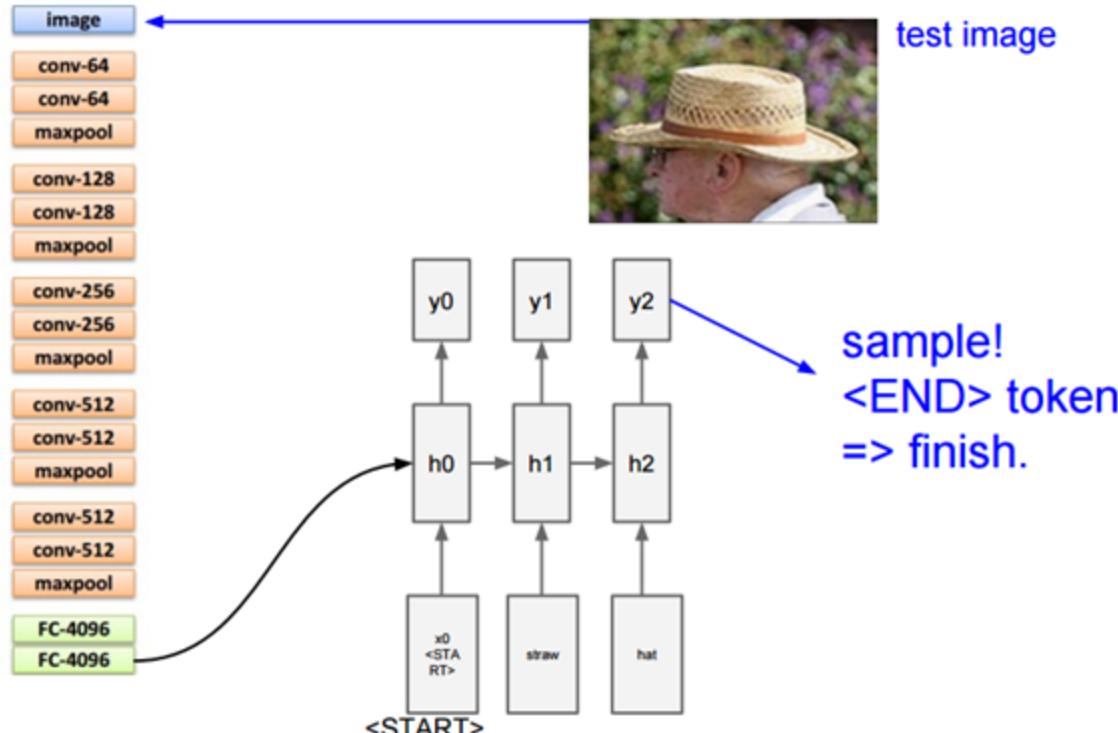
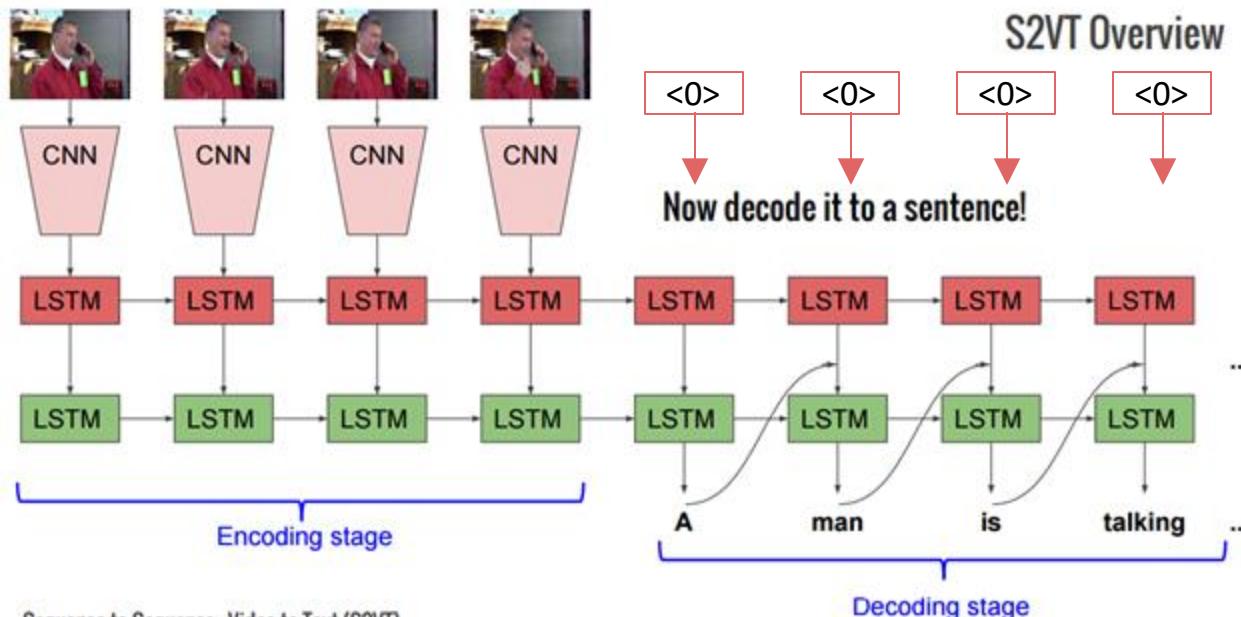


Image Captioning



Video Captioning



Video Captioning

Correct descriptions.



S2VT: A man is doing stunts on his bike.



S2VT: A herd of zebras are walking in a field.



S2VT: A young woman is doing her hair.



S2VT: A man is shooting a gun at a target.

(a)

Relevant but incorrect descriptions.



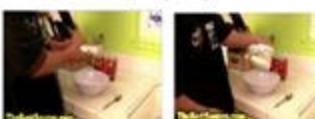
S2VT: A small bus is running into a building.



S2VT: A man is cutting a piece of a pair of a paper.



S2VT: A cat is trying to get a small board.



S2VT: A man is spreading butter on a tortilla.

(b)

Irrelevant descriptions.



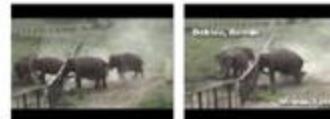
S2VT: A man is pouring liquid in a pan.



S2VT: A polar bear is walking on a hill.



S2VT: A man is doing a pencil.



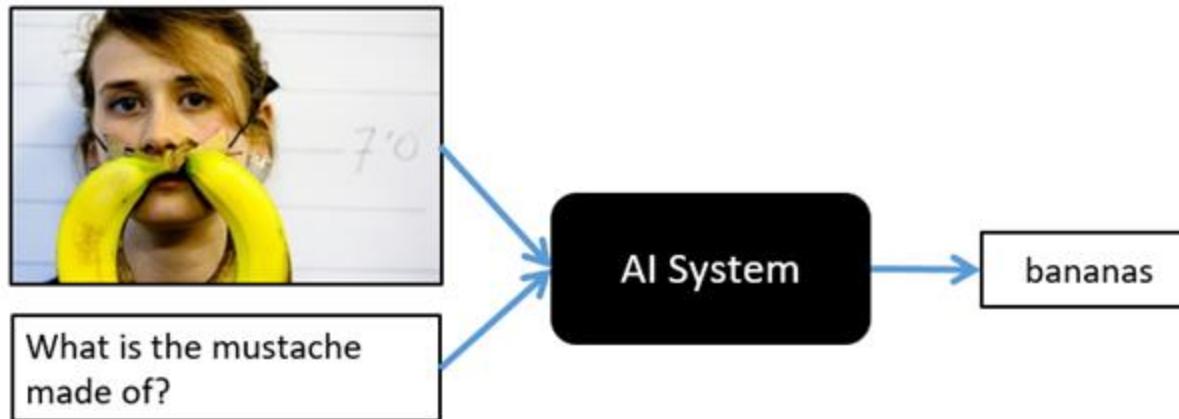
S2VT: A black clip to walking through a path.

(c)

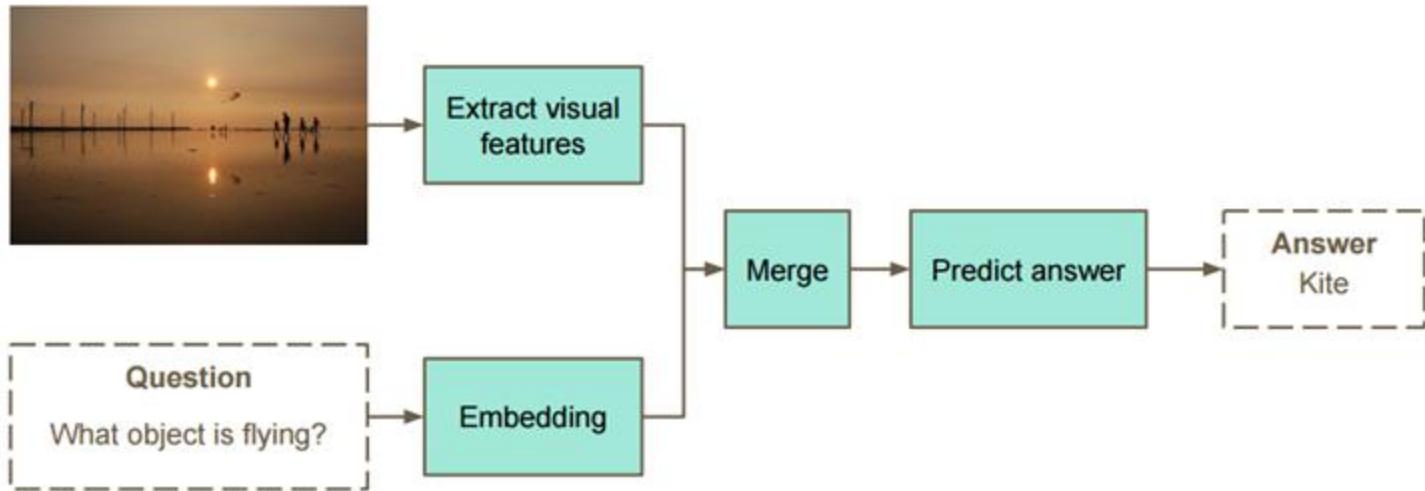
Applications of RNNs in CV

- Image/Video Captioning
- **Visual Question Answering**
- Action / Activity Recognition
from video

Visual Question Answering



Visual Question Answering



Slide credit: Issey Masuda

Visual Question Answering

Questions

Stump a smart robot! Ask a question about this image that a human can answer, but a smart robot probably can't!

We have built
kitchen, beach
Ask a question
(IMPORTANT: T
the question w

Stump a smart robot!
**Ask a question that a human can answer,
but a smart robot probably can't!**

can recognize the scene (i
smart robot!
should not be able to ans
ns below:



- **Do not repeat questions.** Do not ask the same questions or the same questions with minor variations over and over again across images. Think of a new question each time specific to each image.
- Each question should be a **single question**. **Do not ask questions that have multiple parts** or multiple sub-questions in them.
- **Do not ask generic questions** that can be asked of many other images. Ask questions specific to each image.

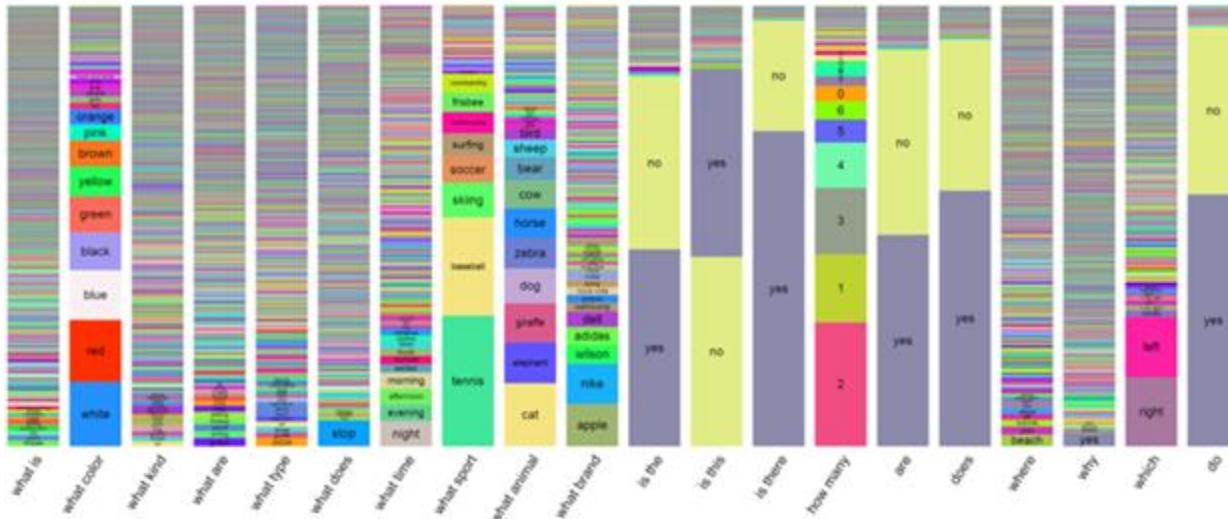
Please ask a question about this image that a human can answer "if" looking at the image (and not otherwise), but would stump this smart robot:

Q1: Write your question here to stump this smart robot.

Slide credit: Devi Parikh

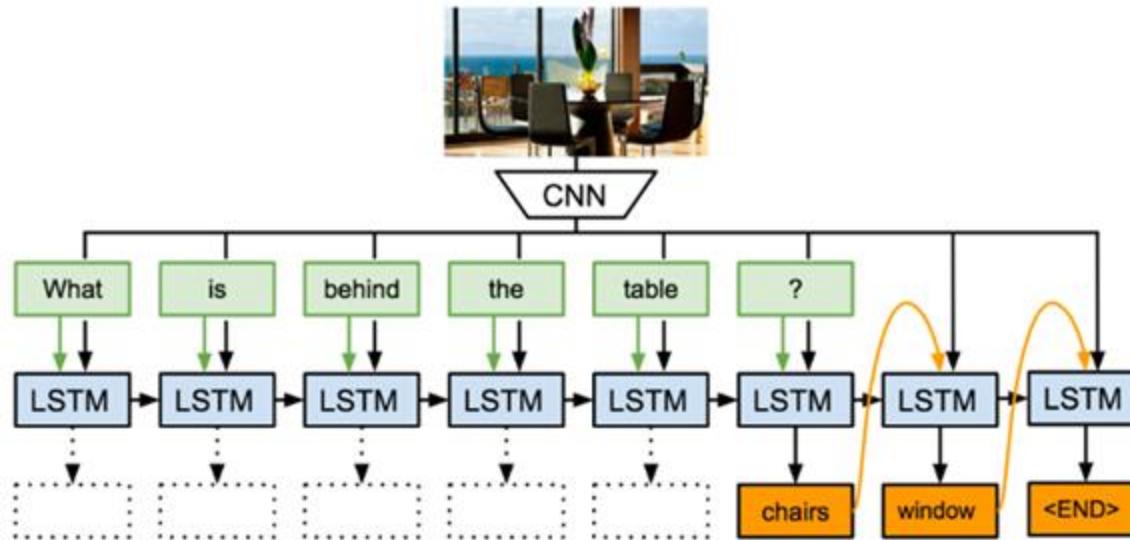
Visual Question Answering

Answers



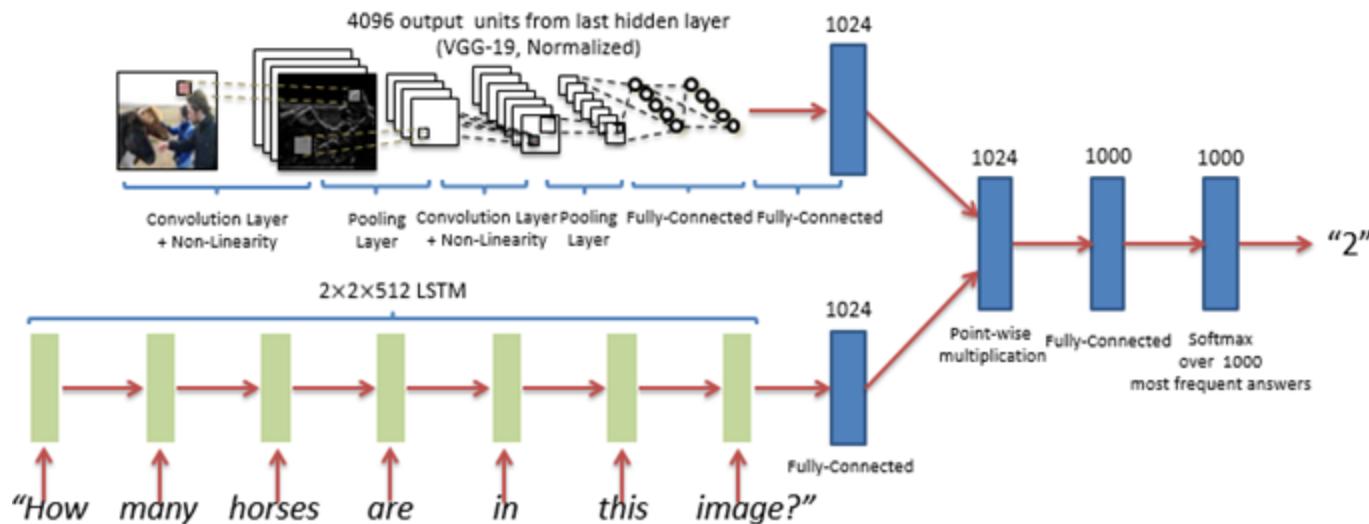
Visual Question Answering

[Malinowski 2015]



Visual Question Answering

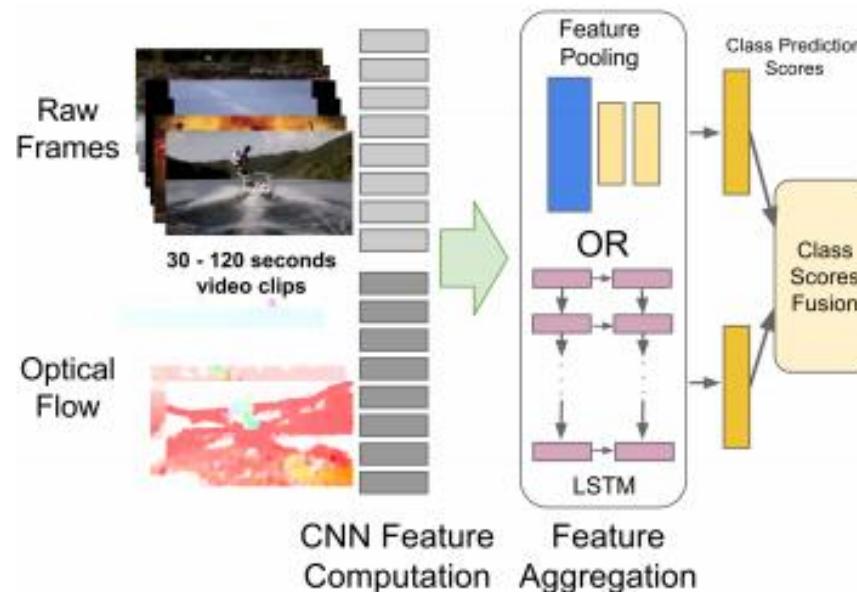
[Lu 2015]



Applications of RNNs in CV

- Image Captioning
- Visual Question Answering
- **Action / Activity Recognition
from video**

Action/Activity Recognition



Action/Activity Recognition

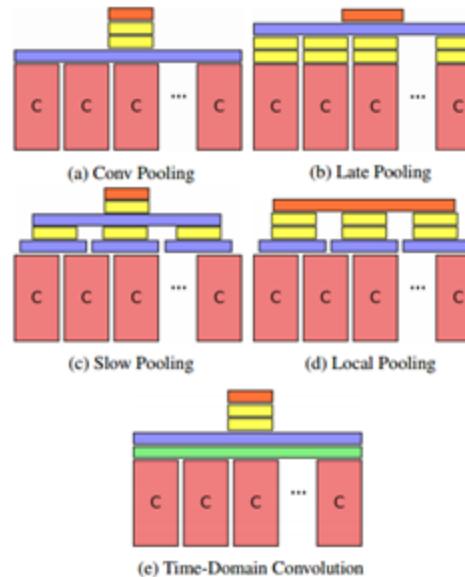
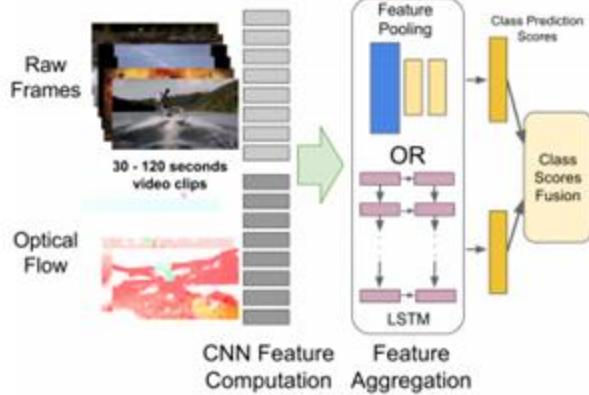
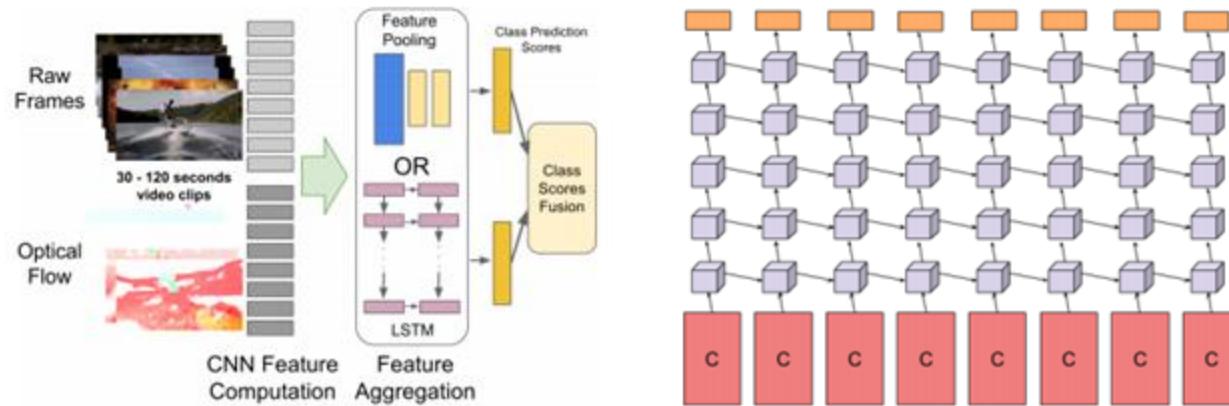


Figure 2: Different Feature-Pooling Architectures: The stacked convolutional layers are denoted by "C". Blue green, yellow and orange rectangles represent max-pooling time-domain convolutional, fully-connected and softmax layers respectively.

Action/Activity Recognition



Action/Activity Recognition

Method	Clip Hit@1	Hit@1	Hit@5
Conv Pooling	68.7	71.1	89.3
Late Pooling	65.1	67.5	87.2
Slow Pooling	67.1	69.7	88.4
Local Pooling	68.1	70.4	88.9
Time-Domain Convolution	64.2	67.2	87.2

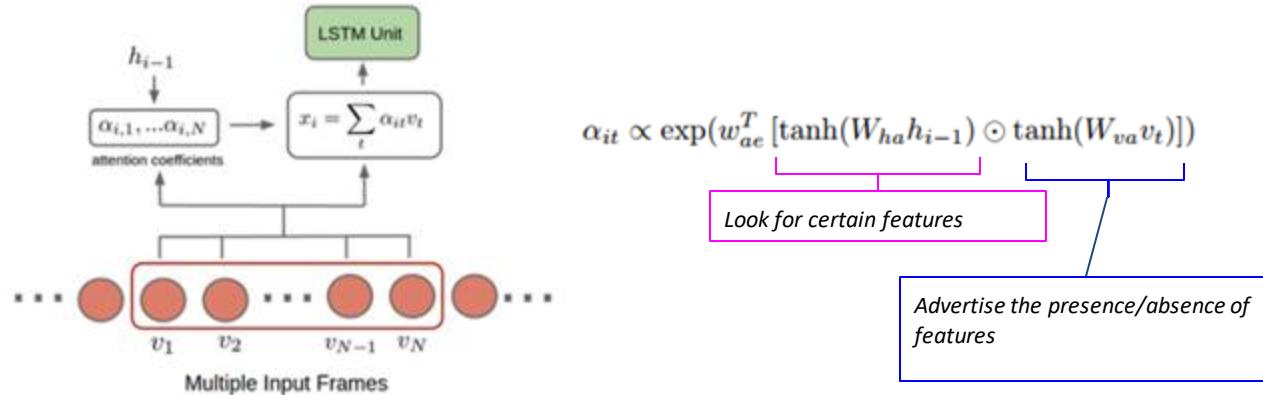
Method	Hit@1	Hit@5
AlexNet single frame	63.6	84.7
GoogLeNet single frame	64.9	86.6
LSTM + AlexNet (fc)	62.7	83.6
LSTM + GoogLeNet (fc)	67.5	87.1
Conv pooling + AlexNet	70.4	89.0
Conv pooling + GoogLeNet	71.7	90.4

Take-home : Consider pooling schemes as an alternative to RNNs

Action/Activity Recognition



Action/Activity Recognition

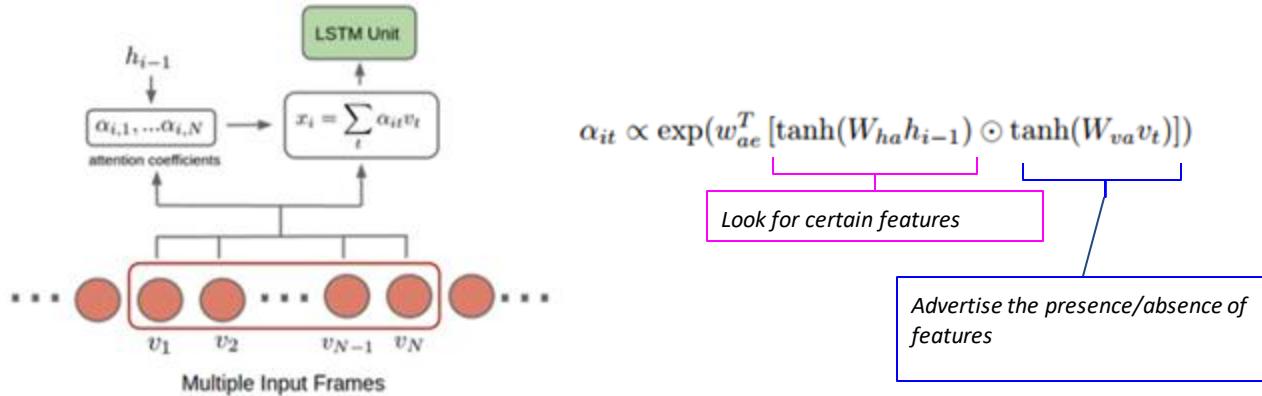


$$y_t = \sum_i \beta_{it} p_{it}$$

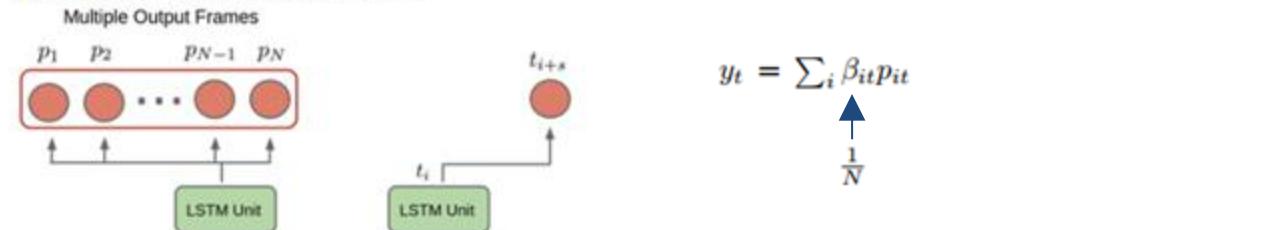
\uparrow
 $\frac{1}{N}$

Every Moment Counts: Dense Detailed Labeling of Actions in Complex Videos, Yeung et al., 2015

Action/Activity Recognition

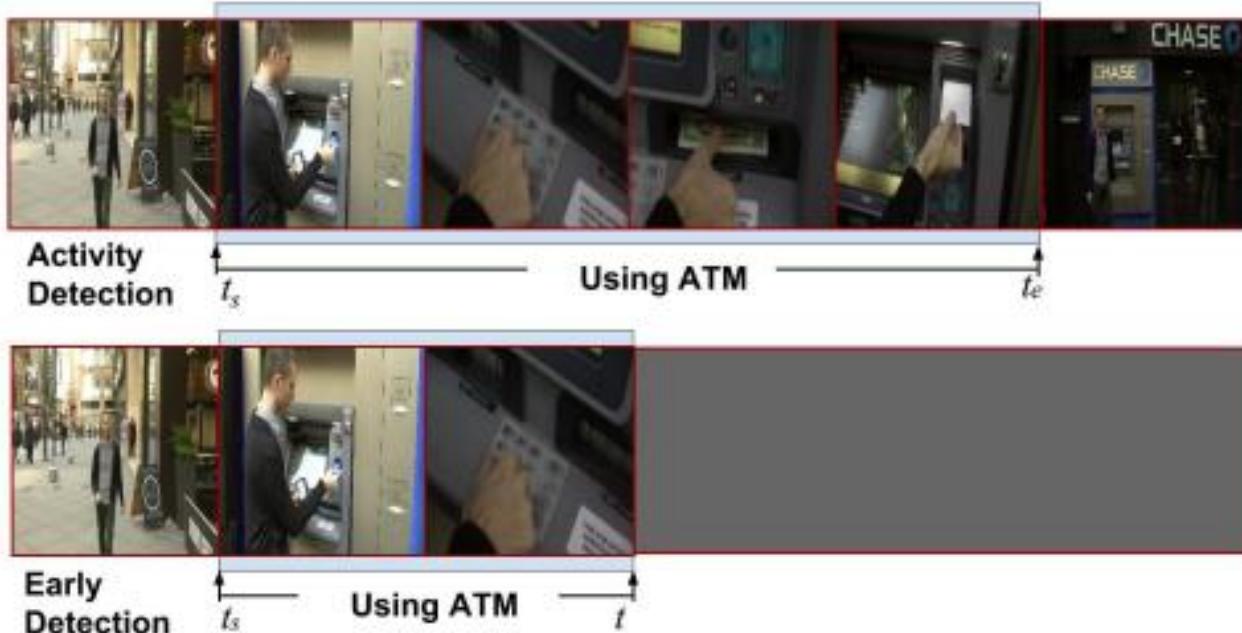


(a) Connections to multiple inputs.

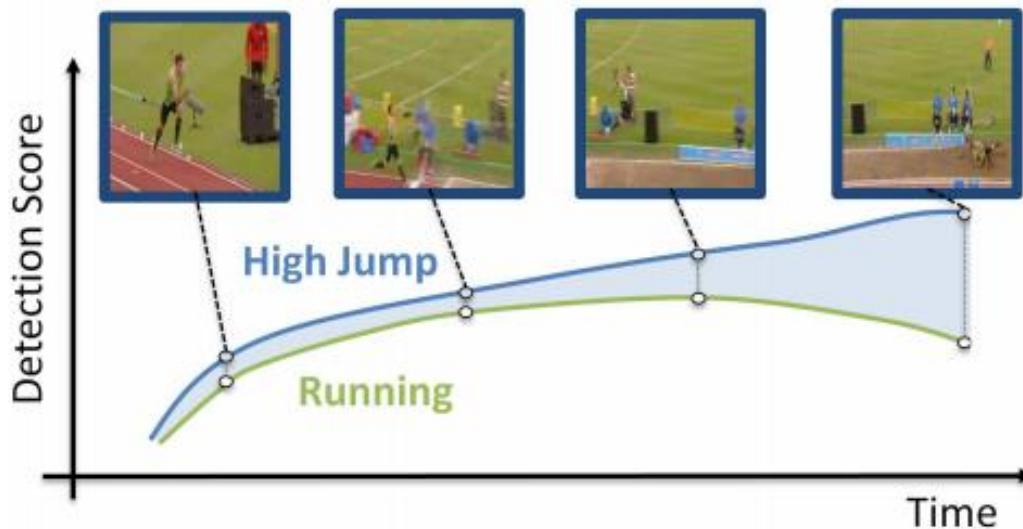


(b) Multiple outputs. (c) Variant: output offset.

Action/Activity Recognition



Action/Activity Recognition



- Score for correct activity should increase over time
- Score margin between correct and incorrect should increase over time

Action/Activity Recognition

- Score for correct activity should increase over time (s-loss)
- Score margin between correct and incorrect should increase over time (m-loss)

In other words, if there is no activity transition at time t , i.e., $y_t = y_{t-1}$, then we want the current detection score to be no less than any previous detection score for the same activity, computing the ranking loss as:

$$\mathcal{L}_s^t = \max(0, p_t^{*y_t} - p_t^{y_t}). \quad (7)$$

On the other hand, if an activity transition happens at time t , i.e., $y_t \neq y_{t-1}$, we want the detection score of the previous activity to drop to zero at t and compute the ranking loss as:

$$\mathcal{L}_s^t = p_t^{y_{t-1}}. \quad (8)$$

$$m_t^y = p_t^y - \max\{p_t^{y'} \mid \forall y' \in \mathcal{Y}, y' \neq y\}$$

In other words, when there is no activity transition at t , we want the current discriminative margin to be no less than any previous margin in the same activity, computing the ranking loss as:

$$\mathcal{L}_m^t = \max(0, m_t^{*y_t} - m_t^{y_t}). \quad (14)$$

If an activity transition happens at time t , we want the discriminative margin of the previous activity to drop and compute the ranking loss as:

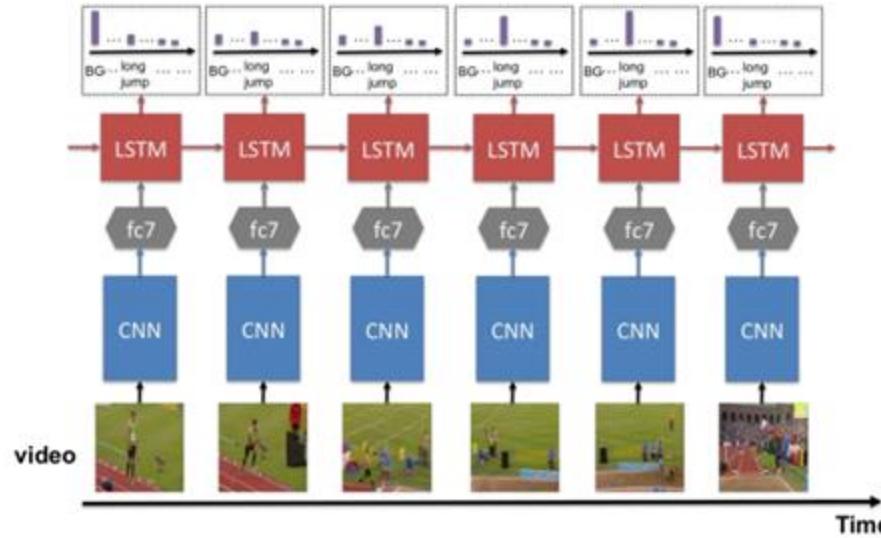
$$\mathcal{L}_m^t = m_t^{y_{t-1}}. \quad (15)$$

$$\mathcal{L}^t = \mathcal{L}_c^t + \lambda_r \mathcal{L}_r^t,$$



$$\mathcal{L}_c^t = -\log p_t^{y_t}$$

Action/Activity Recognition



$$\mathcal{L}^t = \mathcal{L}_c^t + \lambda_r \mathcal{L}_r^t,$$



$$\mathcal{L}_c^t = -\log p_t^{y_t}$$

RNN and CNN

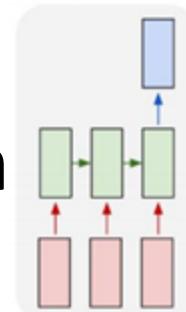
- RNNs = A family of neural networks for processing sequential data
- CNN and RNN are both specialized architectures
 - CNN : processes a grid of values such as an image
 - RNN : processes sequence of values $x(1), \dots, x(\tau)$
 - CNN :

Sharing parameters is key to RNNs

- With separate parameters for each value of time we could not
 - generalize to sequence lengths not seen during training
 - share statistical strength across different sequence lengths and across different positions in time
- (c.f shared weights of conv filter in CNN)

Sharing parameters is key to RNNs

- Compare RNN / CNN weights
 - CNN : Best param setting as determined by independent (image) instances
 - RNN : Best param setting as determined by independent sequences
 - RNN weights have to make sense for current ‘and past’ sequence members



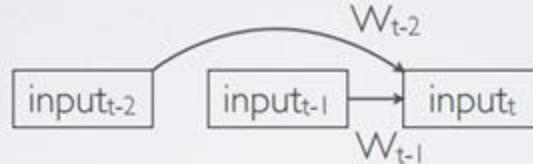
Sharing parameters is key to RNNs

- Sharing is important when a particular info can occur at multiple positions in the sequence
 - Given “I went to Paris in 2012” and “In 2012, I went to Paris”, extract year.
 - Should extract 2012 whether in position 6 or 2 (many → one)
 - A feed-forward net that processes sentences of fixed length would have to learn all of the rules of language separately at each position
 - RNN shares the same weights across several time steps, so can handle this in a scalable manner.

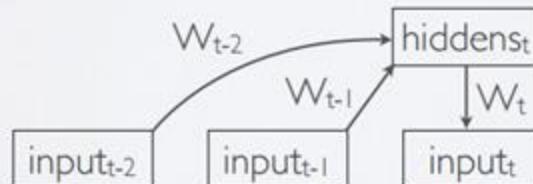
RNNs = Just fancy HMMs ?

- Memoryless models for sequences:

- ▶ **Autoregressive models:** Predict the next input in a sequence from a fixed number of previous inputs using “delay taps”.

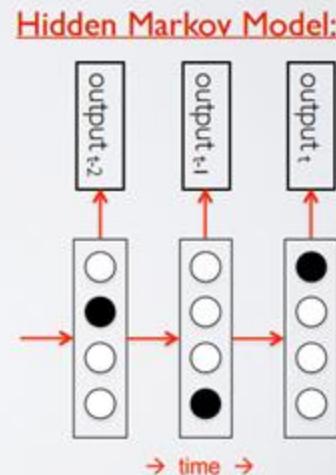


- ▶ **Feed-forward neural networks:** Generalize autoregressive models by using non-linear hidden layers.



RNNs = Just fancy HMMs ?

- If the model could possess a dynamic hidden state:
 - Can store long term information.
- Dynamic Bayes networks (DBNs) can possess internal dynamic state
 - Example 1: Linear Dynamic System with Gaussian noise model (Kalman Filter)
 - Example 2: Discrete state, arbitrary observation type (Hidden Markov Model)
 - State is **not observed**, must be **inferred**.
 - Represent probability across **N** states with **N** numbers.
 - Efficient algorithms exist for HMM inference.



RNNs = Just fancy HMMs ?

LIMITATIONS OF HMMS

- Consider modelling sentence production with an HMM.
- Generation procedure:
 - At each time step t , sample state s_t given state s_{t-1} .
 - Everything important about the past outputs (output 0, ..., output $t-1$) is must be summarized in this choice of state
 - So with N states, it can only remember $\log(N)$ bits of the past.

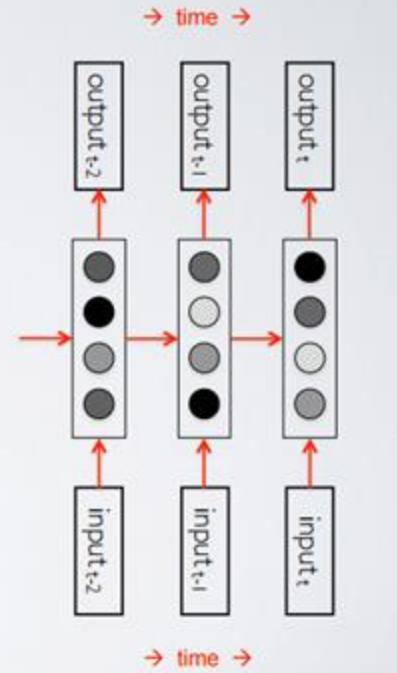
RNNs = Just fancy HMMs ?

LIMITATIONS OF HMMS

- Consider modelling sentence production with an HMM.
- Generation procedure:
 - At each time step t , sample state t given state $t-1$.
 - Everything important about the past outputs (output 0, ..., output $t-1$) is must be summarized in this choice of state
 - So with N states, it can only remember $\log(N)$ bits of the past.
- Consider trying to generate the second half of the sentence with the first have already generated.
 - Syntax needs to fit (number and tense agreement).
 - Semantics need to fit.
- These aspects combined could be say 100 bits that need to be conveyed from the first half to the second. 2^{100} is big!

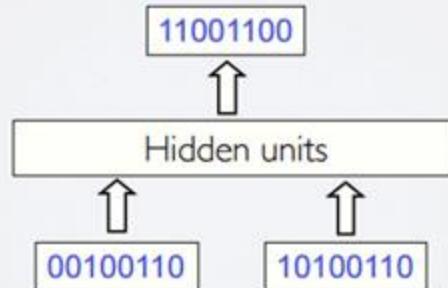
RNNs = Just fancy HMMs ?

- RNNs are very powerful, because they combine two properties:
 - ▶ **Distributed hidden state**: can efficiently store a lot of information about the past.
 - Note: real valued activations not 1-of-N
 - ▶ **Non-linear dynamics**: can update their hidden state in complicated ways.

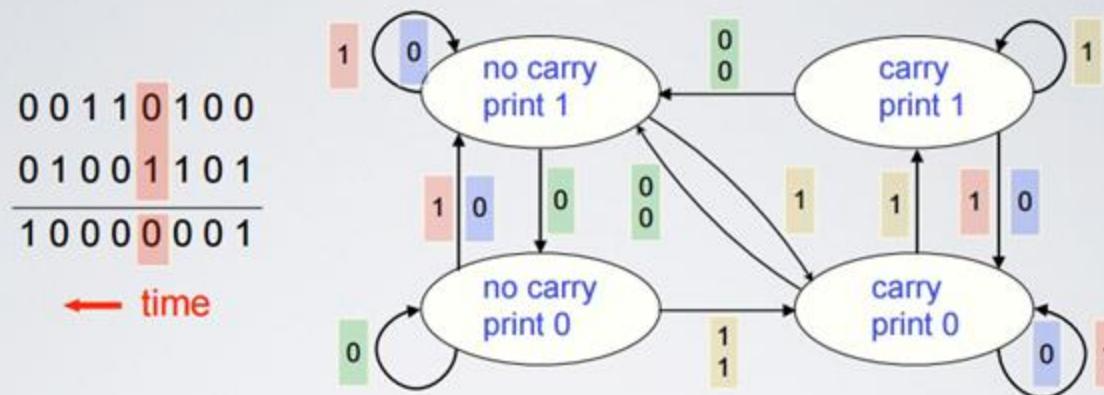


EXAMPLE: BINARY ADDITION

- Can we train a **feed-forward net** to do binary addition?
- **Problems:**
 - We must decide in advance the maximum number of digits in each number.
 - The processing applied to the beginning of a long number does not generalize to the end of the long number because it uses different weights.
- Feed-forward nets do not generalize well on binary addition.



BINARY ADDITION AS A DYNAMIC PROCESS



- Moves from right to left over two input numbers.
- Finite state atomaton
 - Decides what transition to make by looking at the next column.
 - It prints after making the transition.

RNN FOR BINARY ADDITION

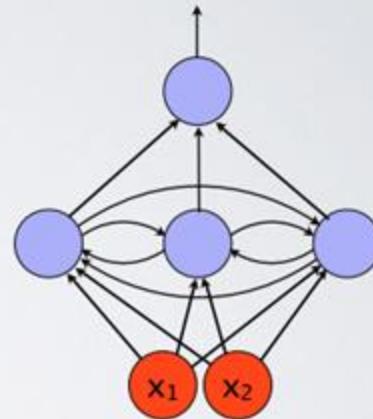
- Network has two input units and one output unit.
 - ▶ Each input corresponds to a digit from a distinct binary number.
- Desired output at each time step is the output for the column that was provided as input two time steps ago.
 - ▶ Takes one time step to update the hidden units based on the two input digits.
 - ▶ Takes another time step for the hidden units to cause the output.

$$\begin{array}{r} 00110100 \\ 01001101 \\ \hline 10000001 \end{array}$$

← time

RNN FOR BINARY ADDITION

- RNN solution has 3 hidden units have all possible interconnections in all directions:
 - Allows hidden activity pattern at one time step to vote for the hidden activity pattern at the next time step.
- Input units have feedforward connections that allow them to vote for the next hidden activity pattern.



Getting targets when modeling sequences

- When applying machine learning to sequences, we often want to turn an input sequence into an output sequence that lives in a different domain.
 - E. g. turn a sequence of sound pressures into a sequence of word identities.*

Getting targets when modeling sequences

- When applying machine learning to sequences, we often want to turn an input sequence into an output sequence that lives in a different domain.
 - *E. g. turn a sequence of sound pressures into a sequence of word identities.*
- When there is no separate target sequence, we can get a teaching signal by trying to predict the next term in the input sequence.
 - **The target output sequence is the input sequence with an advance of 1 step.**
 - This seems much more natural than trying to predict one pixel in an image from the other pixels, or one patch of an image from the rest of the image.
 - For temporal sequences there is a natural order for the predictions.

Getting targets when modeling sequences

- When applying machine learning to sequences, we often want to turn an input sequence into an output sequence that lives in a different domain.
 - *E. g. turn a sequence of sound pressures into a sequence of word identities.*
- When there is no separate target sequence, we can get a teaching signal by trying to predict the next term in the input sequence.
 - The target output sequence is the input sequence with an advance of 1 step.
 - This seems much more natural than trying to predict one pixel in an image from the other pixels, or one patch of an image from the rest of the image.
 - For temporal sequences there is a natural order for the predictions.
- Predicting the next term in a sequence blurs the distinction between supervised and unsupervised learning.
 - It uses methods designed for supervised learning, but it doesn't require a separate teaching signal.